

Argus : Debugging Performance Issues in Modern Applications with Interactive Causal Tracing

Abstract

Prior systems used causal tracing, a powerful technique that traces low-level events and builds dependency graphs, to diagnose performance issues. However, they all assume that precise dependencies can be inferred from low-level tracing by either limiting applications to using only a few supported communication patterns or relying on developers to manually provide dependency schemas upfront for all involved components. Unfortunately, based on our own study and experience of building a causal tracing system for MacOS, we found that it is extremely difficult, if not impossible, to build precise dependency graphs. We report patterns such as data dependency, batch processing, and custom communication primitives that introduce imprecision. We present Argus, a practical system for effectively debugging performance issues in modern desktop applications despite the imprecision of causal tracing. Argus lets a user easily inspect current diagnostics and interactively provide more domain knowledge on demand to counter the inherent imprecision of causal tracing. We implemented Argus in MacOS and evaluated it on 11 real-world, open spinning-cursor issues in widely used applications. The root causes of these issues were largely previously unknown. Our results show that Argus effectively helped us locate all root causes of the issues and incurred only 1% CPU overhead in its system-wide tracing.

1. Introduction

Today’s web and desktop applications are predominantly parallel or distributed, making performance issues in them extremely difficult to diagnose because the handling of an external request is often spread across many threads, processes, and asynchronous contexts instead of in one sequential execution segment [13]. To manually reconstruct this graph of execution segments for debugging, developers have to sift through a massive amount of log entries and potentially code of related application components [6, 24, 20, 15, 22]. More often than not, developers give up and resort to guessing the root cause, producing “fixes” that sometimes make the matter worse. For instance, a bug in the Chrome browser engine causes a spinning cursor in MacOS when a user switches the input method [2]. It was first reported in 2012, and developers attempted to add timeouts to work around the issue. Unfortunately, the bug has remained open for seven years and the timeouts obscured diagnosis further.

Prior work proposed *Causal tracing*, a powerful technique to construct request graphs (semi-)automatically [23]. It does so by inferring (1) the beginning and ending boundaries of the

execution segments (vertices in the graph) involved in handling a request; and (2) the causality between the segments (edges)—how a segment causes others to do additional handling of the request. Prior causal tracing systems all assumed certain programming idioms to automate inference. For instance, if a segment sends a message, signals a condition variable, or posts a task to a work queue, it wakes up additional execution segments, and prior systems assume that wake-ups reflect causality. Similarly, they assume that the execution segment from the beginning of a callback invocation to the end is entirely for handling the request that causes the callback to be installed [23, 16]. Causal tracing is quite effective at aiding developers to understand complex application behaviors and debug real-world performance issues.

Unfortunately, based on our own study and experience of building a causal tracing system for the commercial operating system MacOS, we found that modern applications frequently violate these assumptions. Hence, the request graphs computed by causal tracing are imprecise in several ways. First, an inferred segment may be larger than the actual event handling segment due to batch processing. Specifically, for performance, an application or its underlying frameworks may bundle together work on behalf of multiple requests with no clear distinguishing boundaries. For instance, WindowServer in MacOS sends a reply for a previous request and receives a message for the current request using one system call `mach_msg_overwrite_trap`, presumably to reduce user-kernel crossings.

Second, the graphs may be missing numerous causal edges. For instance, consider ad hoc synchronization [19] via shared-memory flags: a thread may set “`flag = 1`” and wake up another thread waiting on “`while(!flag);`” to do additional work. Even within one thread, the code may set a data variable derived from one request and later uses it in another request (e.g., the buffer that hold the reply in the preceding WindowServer example). Although the number of these flags may be small, they often express critical causality, and not tracing them would lead to many missing edges in the request graph. However, without knowing where the flags reside in memory, a tool would have to trace all memory operations, incurring prohibitive overhead and adding many superfluous edges to the request graph.

Third, in any case, many inferred edges may be superfluous because wake-ups do not necessarily reflect causality. Consider an `unlock()` operation waking up a thread waiting in `lock()`. This wake-up may be just a happens-stance and the developer intent is only mutual exclusion. However, the actual

semantics of the code may also enforce a causal order between the two operations.

We believe that, without detailed understanding of application semantics, request graphs computed by causal tracing are *inherently* imprecise and both over- and under-approximate reality. Although developer annotations can help improve precision [5, 17], modern applications use more and more third-party libraries whose source code is not available. In the case of tech-savvy users debugging performance issues such as a spinning (busy) mouse cursor on her own laptop, the application’s code is often not available. Given the frequent use of custom synchronizations, work queues, and data flags in modern applications, it is hopeless to count on manual annotations to ensure precise capture of request graphs.

In this work, we present Argus, a practical system for effectively debugging performance issues in modern desktop applications despite the imprecision of causal tracing. Argus’s goal is not to construct a per-request causal graph which requires extremely precise causal tracing. Instead, it captures an approximate event graph for a duration of the system execution to aid diagnosis. We designed Argus to be interactive, as a debugger should rightly be, so that its users can easily inspect current diagnostics and guide the next steps of debugging to counter the inherent imprecision of causal tracing. For instance, Argus’s event graph contains many inaccurate edges that represent false dependencies, which the user can address in an interactive manner. When debugging a performance issue using Argus, a user need only make edges relevant to the issue precise. In other words, she can provide schematic information on demand, as opposed to full manual schema upfront for all involved applications and daemons [5].

Moreover, Argus enables users to dynamically control the granularity of tracing using a number of intuitive primitives. The system begins by using always-on, lightweight, system-wide tracing. When a user observes a performance issue (e.g., a spinning cursor), she can inspect the current graph Argus computes, configure Argus to perform finer-grained tracing (e.g., logging call stacks and instruction streams) for events she deems relevant, and trigger the issue again to construct a more detailed graph for diagnosis. Argus also supports interactive search over the event graphs that contain both normal and buggy executions for diagnosis.

We implemented Argus in MacOS, a widely used commercial operating system. MacOS is closed-source, as are its common frameworks and many of its applications. This environment therefore provides a true test of Argus. We address multiple nuances of MacOS that complicate causal tracing, and built a system-wide, low-overhead tracer.

We evaluated Argus on 11 real-world, open spinning-cursor issues in widely used applications such as the Chromium browser engine and MacOS System Preferences, Installer, and Notes. The root causes of all 11 issues were previously unknown to us and, to a large extent, the public. Our results show that Argus is effective: it helped us find all root causes

of issues, including the Chromium issue that remained open for seven years. Argus is also fast: its systems-wide tracing incurs only 1% CPU overhead overall.

This paper makes the following contributions: our conceptual realization that causal tracing is inherently imprecise and that interactive causal tracing is superior than prior work in debugging performance issues in modern applications; our system Argus that performs system-wide tracing in MacOS with little overhead; and our results diagnosing real-world spinning (busy) cursors and finding root causes for performance issues that have remained open for seven years.

This paper is organized as follows. In Section 2, we present an overview of using Argus and a Chromium example. Section ?? describes our approach to identifying semantic dependencies, and Section 4 describes our tracing implementation. In Section 5 we present other case studies, and Section 6 contains our performance evaluation. We summarize related work in Section 7, and end with conclusion in Section 8.

2. Overview

2.1. Argus Work Flow

In this section, we describe the steps a user takes to investigate a performance anomaly with Argus. Figure 1 shows Argus’s work flow, which consists of two phases. A user runs command “Argus start” to enter the system-wide tracing phase, within which Argus logs events as listed in Table 1 (§3.1). Whenever a user detects a performance issue such as a spinning cursor, she runs “Argus debug” to enter the diagnosis phase.

Central to our system is our *event graph*, a generalized control-flow graph which includes inter-thread and inter-process dependencies. Diagnosis and inferences are performed within this graph, in a semi-automated fashion: Argus performs searches and subgraph matching to trace logical events as they flow through the system, and the user can interactively query this graph to understand the problem or provide guidance to Argus. The event graph is described in further detail in §3.1. Next, we describe the graph operations Argus performs leading to a diagnosis.

2.2. Diagnosis with Graph

Algorithm 1 shows how Argus investigates a performance anomaly for users, given a buggy test case and (potentially) a baseline test case for comparison purposes. We assume the user has a program to execute and a buggy test case that can be reproduced eventually; the tracing log collected by Argus would contain a similar event processing normally, or the user can provide a normally executed test case for comparison purpose, listed in line 2. Argus initializes this debugging phase by constructing an event graph from all logged events up to user-specified duration in line 4. The detail of algorithm ComputeGraph is presented in Algorithm 4 (§3.2).

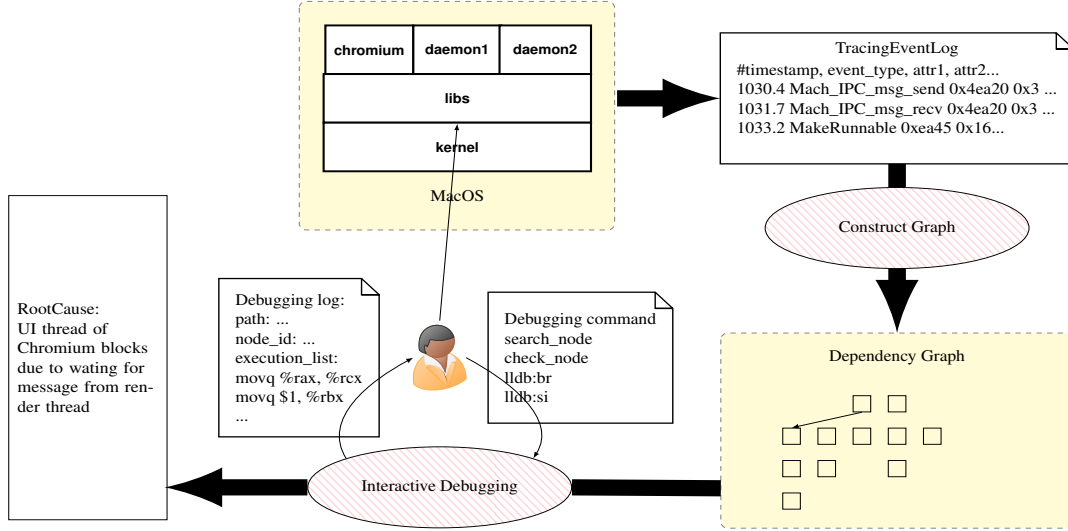


Figure 1: Argus Work Flow

Algorithm 1 Main Argus algorithm.

Require: Program to run + buggy test case + (optional) similar baseline test case

Ensure: Sequence of UI events triggered the performance problem + Nodes from Event Graph

```

1: function ARGUS MAIN
2:   TracingEvents  $\leftarrow$  Run program under Argus, trigger baseline case
   (if possible) and buggy test case
3:   HeuristicsSet  $\leftarrow$  {default heuristics}
4:   ControlFlowGraph  $\leftarrow$  ComputeGraph(HeuristicsSet,
   TracingEvents)
5:   SpinningNode  $\leftarrow$  search ControlFlowGraph with timestamp when
   spinning cursor displays
6:   if SpinningNode == CPUbusyexecution(livelock) then
7:     run backward traversal of edges until UI event found
8:   else
9:      $N_0 \leftarrow$  SpinningNode
10:     $N_1 \leftarrow$  FindSimilarNode(ControlFlowGraph,  $N_0$ )
11:    call AssistedGraphDiff(ControlFlowGraph,  $N_0$ ,  $N_1$ )
12:    if new_hueristics were added then
13:      go to 4
14:    end if
15:  end if
16:  Return set of UI events and root cause related Nodes
17: end function

```

The exact searches and queries performed on the graph depend on the bug under investigation. Consider a common performance bug on MacOS, the *spinning cursor*, which indicates the current application’s main thread has not processed any UI events for over two seconds. Argus queries the event graph to find the ongoing event in the application’s main thread concurrent to the display of the spinning cursor shown in line 5. Upon examining what the main thread is actually doing, the user may encounter three potential cases. First, the thread may be busy performing lengthy CPU operations (which take longer than two seconds). Second, the thread may be in a blocking wait. Third, the main thread may be in a yield loop, which is highly indicative it is waiting on a data flag (e.g., “while(!done)

thread_switch();”). Line 7 to 14 is to track events and synchronization primitives throughout the system to figure out the root cause, in terms of user input events and bug in code, for the three cases.

For the first case, shown in line 7, Argus will examine the stack trace and find a sequence of events that led to the current CPU processing. If more specific tracing is required, the user can rerun the program with more heavyweight instrumentation enabled for any portion of the code’s execution, gathering a precise sequence of calls or even instructions executed.

On the contrary, for the rest cases, line 9 to line 14, Argus locates the node N_1 , either a blocking wait or a yield loop, in our graph when the spinning beachball is triggered. Argus applies Algorithm 2 to locate a node N_0 in our graph that corresponds to the baseline version of the hanging node N_1 .

For the second case of block waiting, Argus takes the assumption there usually exists a chain of blocking thread and the initial blocking thread is the most valuable for debugging. Consequently, Argus diffs the subgraphs between the blocking version and the baseline version to determine their difference with Algorithm 3.

In the third case, the main thread is in a yield loop, which is highly indicative it is waiting on a data flag (e.g., “while(!done) thread_switch();”). To discover a data flag, the user re-runs the application with Argus to collect instruction traces of the concurrent event in both the normal and spinning cases and detects where the control flow diverges. She then reruns the application with Argus to collect register values for the basic blocks before the divergence and uncovers the address of the data flag. She then configures Argus to log accesses to the flag during system-wide tracing. Finally, she can recursively apply Argus to further diagnose “the culprit of the culprit”.

Based on our results and experience, the first case is the most common, but the second and third represent more se-

vere bugs. Long-running CPU operations tend to be more straightforward to diagnose with existing tools. Blocking or yielding cases involve multiple threads and processes, and are extremely hard to understand and fix even for the application's original developers. Therefore, issues remain unaddressed for years and significantly impact the user experience. Algorithm 1 is semi-automated but can integrate user input at each stage to leverage hypotheses or expert knowledge as to why a hang may occur.

Algorithm 2 Argus Find similar node algorithm.

Require: EventTypesSet + SpinningNode + Graph

Ensure: SimilarNodeSet

```

1: function FINDSIMILARNODE
2:    $thread \leftarrow$  thread of  $Node_b$  locking
3:   for  $Node_i$  in  $thread$  other than  $Node_b$  locking do
4:      $iter_i \leftarrow$  iterator of first event in  $Node_i$ 
5:      $iter_b \leftarrow$  iterator of first event in  $Node_b$  locking
6:     while  $iter_b \neq \text{end}$  and  $iter_i \neq \text{end}$  do
7:       while  $iter_i \neq \text{end}$  and  $\text{Event}(iter_i) \notin \text{EventTypesSet}$  do
8:          $iter_i++$ 
9:       end while
10:      while  $iter_b \neq \text{end}$  and  $\text{Event}(iter_b) \notin \text{EventTypesSet}$  do
11:         $iter_b++$ 
12:      end while
13:      if  $iter_b \neq \text{end}$  and  $iter_i \neq \text{end}$  then
14:         $if_{eq} \leftarrow \text{CompareEvent}(\text{Event}(iter_b), \text{Event}(iter_i))$ 
15:        if  $if_{eq} == \text{False}$  then
16:          Break
17:        end if
18:      end if
19:      end while
20:      if  $iter_b == \text{end}$  and  $iter_i == \text{end}$  then
21:        if  $\text{wresult of } Node_i \neq \text{wresult of } Node_b \text{ locking or } \text{timespan of } Node_b \text{ locking} - \text{timespan of } Node_i \geq \text{threshold}$  then
22:          Put Node into  $\text{SimilarNodeSet}$ 
23:        end if
24:      end if
25:    end for
26:  end function

```

Algorithm 3 Argus Assisted graph diff algorithm.

Require: SpinningNode + BaselineNode + Graph

Ensure: Possible root cause of thread blocking

```

1: function ASSISTEDGRAPHDIFF
2:    $\text{ResultSet} \leftarrow \{\}$ 
3:    $\text{SlicingPath} \leftarrow \text{BackwardSlicing on BaseLineNode}$ 
4:    $T_{spinning} \leftarrow$  timestamp of spinning beachball for  $\text{SpinningNode}$ 
5:   for  $thread \leftarrow$  thread of  $Node_i$  in  $\text{SlicingPath}$  do
6:      $T_{normal} \leftarrow$  timestamp of last event in  $Node_i$ 
7:      $Node_b \text{ locking} \leftarrow$  thread blocks during  $(T_{normal}, T_{spinning})$ 
8:     if  $Node_b \text{ locking} \neq \text{NULL}$  then
9:       put  $Node_b \text{ locking}$  into  $\text{ResultSet}$ 
10:    end if
11:    if  $\text{UIEvents}$  in Main UI thread then
12:      put  $\text{UIEvent}$  into  $\text{ResultSet}$ 
13:    end if
14:  end for
15:  Return  $\text{ResultSet}$ 
16: end function

```

2.3. Argus Assisting Algorithm

We mentioned previously in §2.2 that Argus explores the event graph semi-automatically to narrow in on relevant nodes in event graph for diagnosis. Algorithms that require no user intervention are described in this section.

Find Similar Node Algorithm FindSimilarNode in Algorithm 2 is meant to find nodes which have the same high level semantics but different execution results compared to the spinning node. For example, both of them are waiting on locks, but one returns successfully and the other not. Argus defines a subset of events that preserve semantics, which contains System_call, Back_trace, NSApp_event, Mach_message, Wait, Dispatch_enqueue, Dispatch_invoke, Runloop_submit, Runloop_invoke, Share_flag_read and Share_flag_write. The comparison of events depends on event type and its category in Table 1. For connection events, Argus compares their peers. For semantics Events, Argus compares the content of the event attributes, for example, the system call number. The wait results of wait event usually differentiate the normal execution and buggy case. Argus differentiate the cases with wait results by default. Depending on the report detail of the spinning node, user can change the comparing metrics.

Graph Diff Algorithm Argus tries to figure out the root cause for the second case (§2.2) with Algorithm 3. Argus assumes that the blocking in the main thread usually caused by the unresponsive of another thread. In the algorithm, with the input of the blocking node in main thread, Argus carries out a backward path slicing from the node from the baseline case in the graph. The path consists of nodes from multiple threads. Argus traverses every thread and confines the search within the time interval from the base line node in the thread and the timestamp when the blocking node appears in main thread. Among them Argus expects to find a potential blocking in some threads and reports them. At the same time Argus also collects the UIEvents into return set.

2.4. Chromium Spinning Cursor Example

One of the authors experienced first-hand the aforementioned performance issue in Chromium, an open-source browser engine that powers Google Chrome and, starting recently, Microsoft Edge [1]. She tried to type in the Chromium search box a Chinese word using SCIM, the default Chinese Input Method Editor that ships with MacOS. The browser appeared frozen and the spinning cursor occurs for a few seconds. Afterwards everything went back to normal. This issue is reproducible and always ruins her experience, but it is quite challenging to diagnose because two applications Chromium and SCIM and many daemons ran and exchanged messages. This issue was reported by other users for other non-English input methods, too.

To diagnose this issue with Argus, the author started system-wide tracing, and then reproduced the spinning cursor with a Chinese search string typed via SCIM while the page was load-

ing. It produced normal cases for the very first few characters, and the browser got blocked with the rest input as spinning cases. The entire session took roughly five minutes.

She then ran Argus to construct the event graph. The graph had 2,749,628 vertexes and 3,606,657 edges, almost fully connected. It spans across 17 applications; 109 daemons including `fontd`, `mdworker`, `nsurlsessiond` and helper tools by applications; 126 processes; 679 threads, and 829,287 IPC messages. Given the scale of the graph and the diverse communication patterns, it would be extremely challenging for prior automated causal tracing tools [3, 23, 4, 8] because they handle a fairly limited set of patterns. Tools that require manual schema [5, 17], would be prohibitive because developers would have to provide schema for all involved applications and daemons.

Next she ran Argus to find the event in the main thread of the browser process. Argus returned a `cv_timed_wait` event (Figure ??) that blocked the main thread for a few seconds. Inspection of the lightweight call stack revealed that this wait happened within a call to `TextInputClientMac::GetFirstRectForRange`. Without knowing the application’s semantics, she could not understand this method. Thus she ran Argus to compare the spinning case to a normal case. Argus searched in the main thread of the browser process for vertexes similar to this wait waiting vertexes similar to this wait, found three, and confirmed with the user which one she wanted.

What caused the wait in the `renderer` main thread though? She thus continued diagnosis and recursively applied Argus to the wait in `renderer`, and got the wake-up path shown in the figure for this wait. Inspection reveal that the `renderer` requested the browser’s help to render Javascript and was waiting for a reply. At this point, a circular wait formed because the browser was waiting for the `renderer` to return the string bounding box and the `renderer` was waiting for the browser to help render Javascript. This circular wait was broken by a timeout in the browser main thread (the `cv_timed_wait` timeout was 1,500 ms). While the system was able to make progress, the next key press caused the spinning cursor to display for another 1,500 ms. The timeout essentially converted a deadlock into a livelock.

3. Argus Graph Computing

3.1. Event Graph

Argus constructs dependency graphs with the events from tracing logs. Tracing logs contain sequence of events per thread. Each event stands for an execution step in a thread. They are grouped into nodes and IPCs, asynchronous calls and thread wakeups serve as edges; some edges can be inside a single thread.

The events traced in Argus are carefully selected for three main purposes: preserve semantics for the node, identify node boundaries inside a thread, and provide connections

between nodes. We classify them into three categories: semantics events, boundary events and connection events, as listed in Table 1.

Event Type	Event Categories
System_call	Semantics
Back_trace	Semantics
NSApp_event ¹	Semantics
Interrupts	Boundary
Sharetime_maintenance	Boundary
Wait	Semantics, Boundary
Dispatch_invoke	Boundary
RunLoop_invoke	Boundary
Mach_message	Boundary, Connection
Wake_up	Connection
Dispatch_enqueue	Connection
RunLoop_submit	Connection
Share_flag_read	Connection
Share_flag_write	Connection

1. User input events dispatched to the Application.

Table 1: Event Type Categories.

Given the prevalent of multi-threading and multi-processing programs, bugs are much more complicated. The long opening bugs are usually have several threads involve, even across process boundaries. As an example, the always timeout on particular synchronization primitive in one thread usually need to trace back to find the other thread that was responsible for signal the primitive. Compared to the existing debugging tools like lldb and spindump, the dependency graph is useful in that 1) it provides thread relationships all over the system across process boundary and timing boundary and 2) it records execution history for an input event before users capture hangs with their eyes.

3.2. Graph Computing Algorithm

In the section, we describe the algorithm Argus uses to generate an event graph. The algorithm has two main steps: construct nodes with heuristics based on boundary events, and generate edges from the connection events.

A node is a sequence of events derived from the execution of a task in a thread. As shown in Algorithm 4, Argus checks events per thread and applies heuristics when a boundary event is checked. Argus provides five default heuristics. Four of them share the idea of general casual tracing, one is used to work around unknown programming paradigms, and more are expected from users to improve graphs. We discuss them in below.

As listed in (§3.3), event sequences for interrupt processing and kernel maintainance are removed from threads in line 7. The second heuristics in line 9 treats a wait event as an end of a node. An wait event usually indicates a thread switches

to other tasks, but it is not always true considering a thread may park due to low thread priority. One of the examples in MacOS is the pause of worker threads draining a low priority dispatch queue. Dispatch queues are FIFO queues to which an application can submit tasks in the form of block objects. Wait events should not divide the block objects, otherwise it may result in a missing connection. To save the integrity of block objects, Argus keeps a counter `callout_level` to mark if the current event is inside a block object. Only the begin and end of `Dispatch_invoke` are served as the boundary for the node, as is shown in line 14. The heuristics also applies to runloop. Runloop is an event processing loop that used to schedule work and coordinate the receipt of incoming events. The begin and end of the work invocation are served as boundaries. The last default heuristics in line 22 is used to work around the difficulty of exploiting batch processing programming paradigms completely, as listed in follows (§3.3). Argus makes use of mach messages to avoid clustering multiple tasks due to unknown batch processing. It defines `IPC_peer_set` to compute the set of mach message receivers/senders for every node. For every mach message, the algorithm checks whether its peer process exists in `IPC_peer_set`, and adds the message into the current node if the condition is true or the set is empty. Otherwise it adds an ending boundary for current node and constructs a new node beginning with the mach_message event.

Edges connect nodes, both intra-thread and inter-thread, with connection events. Argus walks through each connection event type to apply heuristics as follows. First, the return from a wait operation causally depends on the wake-up operation. An edge is defined from the wake-up event to the first event after the wait returns. Argus also add a weak edges from the wait event to the wake-up event in line 26. Mach message is the core of ipc mechanism implemented in kernel, upon which higher level RPC are built. Argus connects the sender and receiver of a mach message. For messages that expect a reply, Argus also connects the receiver of the original message and the sender of the reply message in line 29. As discussed above, dispatch queue and runloop are popular batch processing programming paradigms in MacOS, Argus connects submissions of a task and executions of the task, which are listed in line 34 and line 36 respectively. Similarly, Argus adds edge from a timer armed event to its triggered event in line 38. Shared flag can either be traced with binary instrument, such as `need_display` flag for CoreAnimation in (§3.3), or with breakpoint watcher command line tool provided by Argus. Edges from the flag written to its read are added from line 40. Since share variables are hard to completely exploit, and the causality can be more complicated than the writer-reader pattern, user interaction is expected.

Finally, a event graph returns and is subject to the improvement with more user input heuristics.

Algorithm 4 Argus Compute Graph algorithm.

Require: Heuristics set + parsed tracing events

Ensure: Control flow graph

```

1: function COMPUTEGRAPH
2:   callout_level  $\leftarrow$  0
3:   IPC_peer_set  $\leftarrow$  {}
4:   current_node  $\leftarrow$  NewNode
5:   for Event: Events in a thread do
6:     switch EventType do
7:       case Interrupt or Timeshare_maintenance
8:         Remove following events before return
9:       case Wait
10:        if callout_level equals 0 then
11:          Add end boundary for current node
12:          current_node  $\leftarrow$  NewNode
13:        end if
14:       case Dispatch_invoke or Runloop_invoke
15:         Divide each callout into a node
16:       case Mach_message
17:         if IPC_peer_set  $\neq$   $\emptyset$  and peer  $\notin$  IPC_peer_set then
18:           Add end boundary for current node
19:           current_node  $\leftarrow$  NewNode
20:         end if
21:         update IPC_peer_set
22:       Other Heuristics
23:     end for
24:   for Event: Connection Events do
25:     switch EventType do
26:       case Wake_up
27:         AddWeakEdge(corresponding wait, wake_up)
28:         AddEdge(wake_up, first event after wait returns)
29:       case Mach_message
30:         AddEdge(sender, receiver)
31:         if needs reply then
32:           AddEdge(receiver, reply sender)
33:         end if
34:       case Dispatch_enqueue
35:         AddEdge(Dispatch_enqueue, Dispatch_invoke)
36:       case Runloop_submit
37:         AddEdge(Runloop_submit, RunLoop_invoke)
38:       case Timer
39:         AddEdge(Timer armed, Timer callout)
40:       case Share_flag_read
41:         AddEdge(Share_flag_write, Share_flag_read)
42:       Other Heuristics
43:     end for
44:   end function

```

3.3. Inherent Inaccuracy

However, to construct an accurate and complete dependency graph is difficult, if not impossible. The graph is inherently inaccurate. That one thread wakes up the other thread does not always stand for a causality between them. In implementation, Argus filters out some of the definitive noise in the following types.

- interrupt processing and kernel sharetime maintenance that take over current thread context.
- timer expiration in the kernel which clears up all the waiting threads on an event source.

In addition to the known definitive noise, there exist false connections and missing dependencies based on our experi-

ence of building a dependency graph with traditional causality tracing. We define them as over connection and under connection respectively.

3.3.1. Over Connections

Over connections usually occur when intra-thread boundaries are missing due to unknown batch processing programming paradigms. We list the example patterns we found below.

Dispatch message batching While traditional causal tracing assumes the entire execution of a callback function is on behalf of one request, we found some daemons implement their service loop inside the callback function and create false dependencies. In the following code snippet from the `fontd` daemon, function `dispatch_execute` is installed as a callback to a work from dispatch queue. It subsequently calls `dispatch_mig_server()` which runs the typical server loop and handles messages from different apps.

To avoid incorrectly linking many irrelevant processes through such batching processing patterns, Argus adopts the aforementioned heuristics to split an execution segment when it observes that the segment sends out messages to two distinct processes. Any application or daemon can implement its own server loop this way, which makes it fundamentally difficult to automatically infer event handling boundaries.

```
worker thread in fontd:      main thread in fontd:
//enqueue a block           //dequeue blocks
block = dispatch_mig_sevice; block = dequeue();
dispatch_async(block);      dispatch_execute(block);

//implementation of dipatch_mig_server
dispatch_mig_server()
for(;;) //batch processing
    mach_msg(send_reply, rcv_request)
    call_back(rcv_request)
    set_reply(send_reply)
```

Batching in event processing Message activities inside a system call are assumed to be related traditionally. However, to presumably save on kernel boundary crossings, WindowServer MacOS system daemon uses a single system call to receive data and send data for an unrelated event from different processes in its event loop. This batch processing artificially makes many events appear dependent. We split the execution segments to maintain the independence of the events.

```
//inside a single thread
while() {
    CGXPostReplyMessage(msg) {
        // send _gOutMsg if it hasn't been sent
        push_out_message(_gOutMsg)
        _gOutMsg = msg
        _gOutMessagePending = 1
    }
    CGXRunOneServicePass() {
        if (_gOutMessagePending)
            mach_msg_overwrite(MSG_SEND | MSG_RECV, _gOutMsg)
        else
            mach_msg(MSG_RECV)
        ... // process received message
    }
}
```

3.3.2. Under Connections

On the other hand, under connections mostly result from missing data dependencies. Data dependencies inter/intra threads are usually hard to fully exploit in the initial pass of graph computing.

Data dependency in event processing The case above also illustrates a causal linkage caused by data dependency within one thread. As the code shows, WindowServer saves the reply message in variable `_gOutMsg` inside function `CGXPostReplyMessage`. When it calls `CGXRunOneServicePass`, it sends out `_gOutMsg` if there is any pending message. Argus uses watch point registers to capture events on `_gOutMsg` and establish the causal link between the handling of the previous request and the send of the reply.

CoreAnimation shared flags A worker thread can set a field `need_display` inside a CoreAnimation object whenever the object needs to be repainted. The main thread iterates over all animation objects and reads this flag, rendering any such object. This shared-memory communication creates a dependency between the main thread and the worker so accesses to these field flags need to be tracked.

```
Worker thread:      Main thread:
//needs to update UI: //traverse all CA objects
obj->need_display = 1 if (obj->need_display == 1)
                    render(obj)
```

Spinning cursor shared flag Whenever the system determines that the main thread has hung for a certain period, and the spinning beach ball should be displayed, a shared-memory flag is set. Access to the flag is controlled via a lock, i.e. the lock is used for mutual exclusion, and does not imply a happens before relationship. Thus, Argus captures accesses to these flags using watch-point registers to add causal edges correctly.

```
NSEvent thread:
CGEventCreateNextEvent() {
    if (sCGEventIsMainThreadSpinning == 0x0)
        if (sCGEventIsDispatchToMainThread == 0x1)
            CFRRunLoopTimerCreateWithHandler{
                if (sCGEventIsDispatchToMainThread == 0x1)
                    sCGEventIsMainThreadSpinning = 0x1
                    CGSConnectionSetSpinning(0x1);
            }
}

Main thread:
Convert1CGEvent(0x1);
if (sCGEventIsMainThreadSpinning == 0x1)
    CGSConnectionSetSpinning(0x0);
    sCGEventIsMainThreadSpinning = 0x0;
    sCGEventIsDispatchedToMainThread = 0x0;
```

In addition to batching processing and data dependency, the spurious edge introduced by mutex lock is also a challenge in debugging. The synchronization on mutex lock reflects one thread wakes up the other thread. The scenario can either be a producer-consumer problem or merely mutual exclusion.

Making the graph completely sound without user interaction is almost impossible given the essential attribute of commercial operating system as a grey box.

3.4. User Interactions

As we mentioned above, to figure out all over connections and under connections before hand is almost impossible. Instead, users can find out the discrepancy in event graph after the initial graph computing. For example, we noticed that two unrelated applications connects to one node in the graph, which leads to the manifestation of kernel thread batch processing on timers. Like other causality tracing approaches, Argus is a general framework for MacOS and tested on limited programming paradigms. Allowing user interaction makes its graph more practical and useful case by case.

Users can gradually inject their knowledge to improve the graph in two ways. They can either add heuristics to the graph computing algorithm in Algorithm 4, or binary instrument images with the APIs provided by Argus to expand the boundary or connection event categories.

4. Implementation

We now discuss how we collect tracing events from both kernel and libraries.

4.1. Event Tracing

Current MacOS systems support a system-wide tracing infrastructure built by Apple XXX ???. By default, the infrastructure temporarily stores events in memory and flushes them to screen or disk when an internal buffer is filled. We extended this infrastructure to support larger-scale tests and avoid filling up the disk with a file-backed ring buffer. Subject to configuration, it allows at most 2GB of data per log, which corresponds to approximately 18,560,187 events (about 5 minute with normal operations).

The default tracing points in MacOS provide too limited information to apply causal tracing. As a result, we both patch source code of kernel XXX ?? and binary instrument libraries to gather more tracing data. In Argus, we patched the kernel with 1193 lines of code, and instrumented the libraries including: `libsystem_kernel.dylib`, `libdispatch.dylib`, `libpthread.dylib`, `CoreFoundation`, `CoreGraphics`, `HIToolbox`, `AppKit` and `QuartzCore`, with our binary instrumentation.

4.2. Instrumentation

Most libraries as well as many of the applications used day-to-day are closed-source in MacOS. To add tracing points to such code, techniques such as library preloading to override individual functions are not applicable on MacOS, as libraries use two-level executable namespaces []. Hence, we implemented a binary instrumentation mechanism that allows developers to add tracing at any location in a binary image.

Like Detour [14], we use static analysis to decide which instrumentation to perform, and then enact this instrumentation at runtime. Firstly, users find a location of interest in the image related to a specific event by searching a sequence of instructions. Then the users replace a call instruction to invokes a trampoline target function, in which we overwrite the victimed instructions and produce tracing data with API from Apple. All of the trampoline functions are grouped into a new image, as well as an initialization function which carries out the drop-in replacement. Then command tools from Argus helps to configure the image with the following steps: (1)re-export all symbols from the original image so that the original code can be called Like an shared library; (2)replace the original image with the new one by renaming them to ensure the modifications are properly loaded; (3)invoke the initialization function externally through `dispatch_once` during the loading.

4.3. User Interaction

Tracing Custom Primitives As described in (§3.4), under-connection due to the missing share data dependency requires users' interaction. Argus provides a command line tool which sets the watch point registers to record `share_flag_write` and `share_flag_read` events in ad-hoc manner. This is one way users can easily collect the tracing events, and the more tech-savvy users can also instrument the binary. The tool takes the process id, path to image where the variable is defined and the symbol of the variable as input. We show the simple example how a user ask Argus to trace `_gOutMsgPending` in the following command.

```
./bp_watch pidofWindowServer \  
Path/to/CoreGraphics _gOutMsgPending
```

Argus hooks the watch point break handler in `CoreFoundation` to make sure that it is loaded correctly into the address space of our target application. The handler invokes the event tracing API from Apple to record the value of the shared variable and the operation type: read or write.

Capturing Instructions for Diagnosis Since the output of Argus is the nodes and user input events suspected to cause the busy spinning, more detail may required to verify and fix the bug. After the offline analysis on the graph, Argus provides tools for user to exact the backtrace events from the output and generate a script for conditional debugging.

The debugging scripts go through the instructions of apps and frameworks step by step to capture the parameters tainted by user inputs. At each beginning of a function call, the script records a full callstack for it. Considering the overhead and usefulness, it steps over and only record the return value of APIs from libraries with a filename extension `.dylib`.

The supplementary information are subject to the users review to pinpoint the root cause of spinning beachball on MacOS.

4.4. Limitations

Argus is designed to support interactive debugging of performance issues. To incrementally obtain more fine-grained event traces, it needs to rerun an application to reproduce a performance issue. Thus, if the issue is difficult to reproduce, we have to rely on the log collected by the lightweight system-wide tracing for debugging, and lose the benefits of interactivity. Fortunately, a performance issue that almost never reproduces is probably not as annoying as one that occurs frequently.

We implemented Argus in the closed-source MacOS which presents a harsh test for Argus, but we have not ported Argus to other operating systems yet. It is possible that the ideas and techniques do not generalize to other operating systems. However, modern operating systems share many similarities, and good ideas tend to flow both ways, so we are hopeful that the ideas in Argus are generally applicable. Similarly, the applications and performance issues used in our evaluation may be non-representative.

5. Case Studies

We applied our tool to figure out the design of spinning wait cursor in MacOS, and uncovered root causes of spinning for 11 applications in total. As we have illustrated the Chromium case in Section 2, we will not present it here again.

The spinning beachball is a painful sight for Mac users, signifying that the application is non-responsive. It usually remains for minutes at a time, leaving the user at a loss and unable to do anything productive.

Argus sheds light on the design of the spinning wait cursor with its backward path slicing. We begin the path from the node where spindump, a hang reporting tool, is launched, considering the spindump shares the same triggering condition. For our tracing data, we found out that Spindump is launched after receiving a message from WindowServer which receives a message from the NSEvent thread of the target application first.

We further added call stacks for the messages and revealed two shared variables, “is_mainthread_spinning” and “dispatch_to_mainthread”, are critical in the design. The NSEvent thread of the target application fetches CoreGraphics events from WindowServer, converts and creates NSEvents for the main thread. If the main thread is not spinning with “is_mainthread_spinning” equals 0, “dispatch_to_mainthread” will be set to 1 and a timer is armed. If the main thread processes the next event before the timer fires, nothing happens and the timer gets re-armed. Otherwise, NSEvent thread sends a message inside “CGSConnectionSetSpinning” to WindowServer from the timer handler, and WindowServer notifies the CoreGraphics to draw a spinning wait cursor over the application window.

5.1. System Preferences

System Preferences provides a central location in MacOS to customize system settings. The “Display” pane allows users to configure additional monitors, mirroring or extending their workspace. However, the software does not support disabling monitors that are online (the user must physically unplug monitors). There is a tool called DisableMonitor [10], distributed via GitHub, which addresses this functionality. Surprisingly, we find performance bug in System Preferences when we disable an external monitor, and rearrange the windows in the Display panel afterward. The System Preferences windows freezes for seconds in this situation.

To diagnose this issue, we run the System Preferences app with Argus. We rearrange the displays with two active monitors, and repeat the process with one of them deactivated. Argus collects 132MB data, and constructs dependency graph system-wide in the period, which contains 428,785 vertexes and 320,554 edges.

To diagnose the root cause, Argus first finds out when the NSEvent thread in System Preferences notifies WindowServer to draw the spinning cursor and marks it as t . The node in the main UI thread causing the non-responsiveness must overlap the time interval $(t - 2s, t)$. It is not hard for Argus to tell the spinning node in the main UI thread, either busy processing or blocking.

The spinning node in the main UI thread is dominated by *mach_msg* and *thread_switch*, both of which are meant to wait for available data ping from WindowServer. Noticing the timeouts of *thread_switch*, Argus classifies this case into the third category (§2.2) and heads to find a comparable normal node in the same thread. As its semantics are not descriptive enough to identify its comparable node, Argus extends the comparison with its proceeding nodes.

In the normal case, System Preferences gets rid of the *thread_switch* quickly after receiving messages from WindowServer. It proceeds to *displayReconfigured*. On the other hand, the spinning node ends up sending message for the available datagram ping with *CGSSnarfAndDispatchDatagrams*. By checking the lightweight callstacks, Argus figures out the messages from WindowServer in the previous nodes are responses to data available pings from *activeDisplayNotificationHandler*.

Given the information, we launched the interactive debugging by feeding the debugging script with those APIs mentioned above. We set the method *activeDisplayNotificationHandler* as a breakpoint where the script begins debugging. *displayReconfigured* and *CGSSnarfAndDispatchDatagrams* are used to indicate the end of debugging for the normal case and spinning case respectively.

By diff’ing the two logs, we notice the different branches in *display_notify_proc* called by *activeDisplayNotificationHandler*. The handler depends on received datagram and two shared variables “_gCGWillReconfigureSeen” and

“_gCGDidReconfigureSeen” to finish a display configuration. In both case the first variable is set to indicate the begin of display configuration. In normal case, it receives a datagram which set the variable “_gCGDidReconfigureSeen” in *display_notify_proc* and finish display reconfiguration, while in the spinning case such datagram is never received. Instead, an alternative datagram just drive it through *display_notify_proc* without setting any variable, which causes the repeating *thread_switch* in the handler.

[In conclusion, the bug is... would have to be fixed by... we provide a binary fix...]

5.2. Sequel Pro

Sequel Pro is a fast, easy-to-use Mac database management application for working with MySQL databases. It allows user to connect to database with a standard way, socket or ssh. We noticed spinning beachballs while the network connection was lost and Sequel Pro tried reconnections. In the worst case, work in other tabs get lost if the reconnection fails.

We simulate the unstable network environment by randomly disabling the network for 1 second, and collect data in two cases: 1) network reconnection succeeds quickly, and 2) Sequel Pro loses its connection for a while. 2G tracing data are stored and the graph we generated contains 3,650,832 edges and 2,412,236 vertexes.

Argus starts analysis from the spinning node in NSEvent thread of Sequel Pro and finds the corresponding node in its main UI thread. We notice the execution sequence,

poll, write, wakeup, re-arm timer, read

in the node. Sequel Pro gets blocked on the read. As a result, Argus goes to search the similar normal node, which has the same sequence but is not blocked for long. From the node, Argus goes along the weak edge to figure out which node wakes it up and slice path backwards. Argus carries out the comparison of the nodes in the thread before the normal node gets waken and before the spinning beachball shows up. The close examination tells that the main thread is waiting for the kernel thread, which in turn waits for the ssh thread. Such result could hardly figure out by lldb or spindump as both of them only focus on the callstack snapshot for each process, instead of providing the connections based the execution history system wide.

5.3. TeXstudio

TeXstudio is an integrated writing environment for creating LaTeX documents. In their github we noticed some one reported while he was modifying his bib file, the TeXstudio hanged for minutes. Although the issue was closed by author for incomplete information for reproduction, we reproduce it in our machine. With a relatively large bib file from a LaTeX project opened in a tab, once we touch the bib file through other editors, vim for exmaple, we would have a spinning beachball on the window.

While the beachball is spinning over the application window, the main thread is busy with *QDocumentPrivate::indexOf(QDocumentLineHandle const*, int)*. Slicing the path from the busy node, Argus tells the busy processing is invoked by the callback from daemon fsevents. The advantage of Argus over other debugging tools is it helps to narrow down the root cause with the path. If the user reported the bug with Argus’ report, it should have provided the author more hints to reproduce the bug successfully.

Another operation triggering spinning beachball in TeXstudio is pasting a large amount of context in any file. The spinning node is busy with *QEditor::insertFromMimeData(QMimeData const*)*, which always invokes *match(QNFAMatchContext*, QChar const*, int, QNFAMatchNotifier)*. Path slicing by Argus attributes it to the user’s keyboard inputs of cmd+v. The problematic code resides where developers copy data from the pasteboard.

5.4. Installer

As we were told, any spinning beachball in MacOS indicated a bug. In some situation, the “bugs” confuse users. Installer is one of the case. When Installer pops up a window for privileged permission during the installation of *jdk-7u80-macosx-x64*, moving the cursor out of the popup window triggers spinning beachball.

Argus collected 550M data for Installer, and generated a graph with 392199 edges and 294856 vertexes. It figure out the main thread is blocked in the API *[IFRunnerProxy requestKeyForRights:askUser:]*, which sent a synchronous XPC to daemon *authd* with *xpc_connection_send_message_with_reply_sync*. In conclusion, if the user does not focus on the popup window, it will be treated as a main thread non-responsive bug.

5.5. Notes

Notes is a notetaking app developed by Apple. We encounter spinning beachballs in the app on a relatively long note. Although it works well when the note is generated and stored, the spinning beachball appears when we try to launch Notes the next time. It prevent us from accessing the note for over one hour.

With the collected tracing data, Argus reports Notes is busy with *NSDetectScrollDevicesThenInvokeOnMainQueue*.

5.6. TextEdit

Select, copy, paste, delete, insert and save are common operations for Mac users who do text editing. However, these operations on a large amount of context usually trigger spinning cursors on text editing softwares, including TextEdit developed by Apple. We traced data from six popular text editing softwares, including TextEdit, Microsoft Word, SublimeText, TextMate, CotEditor and TeXstudio. Argus analyzed the tracing data and reported the spinning node for users to inspect.

All of the cases are fell into the non-blocking categories. Our lightweight callstack tells the root cause directly.

For the selection in TextEdit, Argus reports the main thread is busy processing *[NSBigMutableString getCharacters:range:]*, *CFStringGetRangeOfCharacterClusterAtIndex*, *_CFStringInlineBufferGetComposedRange* and *NSFastFillAllGlyphHolesForCharacterRange*. All of them are related to the storage of characters.

For the operation of copy, Argus reports that TextEdit is busy with *get_vImage_converter* and *get_full_conversion_code_fragment* in its main thread. Both of them are called by *[NSTextView(NSPasteboard) _writeRTFDInRanges:toPasteboard:]*.

In the case of paste in TextEdit, it is busy with *DDLookupTableRefLookupCurrentWord_block_invoke_2*, *_RTFGetToken* and *platform_memmove*. All of them called by *-[NSTextView(NSPasteboard) _readSelectionFromPasteboard:types:]*

5.7. Microsoft Word

Copying context in Microsoft Word produced 1.14GB trace file, which contained 474,178 edges and 306,171 vertexes. While the spinning cursor showed up for the application, the main thread was busy with system calls, *lseek*, *fstat64*, *fcntl* and *read* on the same file, *_CFStringCreateImmutableFunnel3* and *platform_memmove* over 2 seconds. They are on the behalf of *-[NSPasteboard _setData:forType:index:usesPboardTypes:]*.

We collected 1.04GB tracing data for Paste operation. Its graph contains 161,921 edges and 110,680 vertexes. The spinning node are dominated by *lseek*, *fstat64*, *fcntl* and *write*.

5.8. SublimeText

We collected 1.16GB trace data for copy in SublimeText. It generates 173,071 edges and 127,488 vertexes. While the cursor is spinning, the main thread of SublimeText is busy encoding characters with *_CFToUTF8Len*. They are called from *CFPasteboardSetData*, which invoked by *px_copy_to_clipboard* in SublimeText source code.

Paste operation in SublimeText also results in the unresponsiveness of UI thread. The size of the tracing data is 505MB. The graph consists of 760,003 edges and 535,106 vertexes. The main thread is busy processing *decode_utf8*, *convert_utf8_to_utf32*, *string_append* and *platform_memmove*. All are called from *px_copy_from_clipboard*.

Deleting large context in SublimeText also causes spinning beachball. Tracing file for it is 447MB, which contains 789,400 edges and 626,629 vertexes. The main thread is busy processing *TokenStorage::substr*.

5.9. TextMate

Spinning cursor appears when a large amount of context get pasted into a file opened with TextMate. We traced 62MB data and got 73,679 edges and 50,860 vertexes. The main thread

is busy with *-[OakTextView paste:]* which iterates through paragraphs, fetchs characters and processed with *CFAttributedString* and *TASCIIEncoder::Encode*.

5.10. CotEditor

Different from other text editing softwares, CotEditor does not trigger spinning beachball while copying a large amount of context. Although paste operation causes the spinning, the path slicing attributes it to the insertion of a new line. The main thread is busy on processing *-[NSBigMutableString characterAtIndex:]* and *CFStorageGetValueAtIndex* in the callstacks. They are invoked after the user hits the return key to insert a line.

Most of the text editing cases fall into the category of thread busy processing. Although sampling with lldb can also figure out the most frequently called APIs, Argus provides the path slicing over the execution history and can connect them to the application events, such as user inputs, daemon notifications and so on.

6. Performance Evaluation

We deployed Argus on a Mac OS X x86 system, model MacBookPro9,2. The model has the Intel Core i5-3210M CPU with 2 cores and 4 thread, 10GB DDR3 memory, and we replaced the hard disk with a 1T SSD. The tracing tool is running in the background 24X7. Once the spinning cursor appears on the screen, we store the tracing data for root cause analysis.

In previous section, we studied real-world software, including TextEdit, Notes, Installer, System Preferences and Chromium. TextEdit, Notes, Installer and System Preferences are distributed by Apple. Our tool can take over most burden searching and comparing work from the user, and make the diagnosis much easier in the wild.

In this section we present the performance impact of the live deployment of Argus. As we use the ring buffer to collect events, the storage cost can be adjusted and is fixed to 2G in our experiments. Since the internal memory used to collect data is fixed to 512M, so the overhead of the memory is pretty low with regards to the memory usage of morden applications. We show the CPU overhead of Argus first for the iBench Running on the system with the tracing on and off. In the following description, we call the environment without the tracing on as bare run, and otherwise tracing run.

	1st	2nd	3rd	4th	5th
bare run	5.98	6.23	6.18	6.05	6.28
tracing run	6.29	6.01	6.09	6.28	6.01

Table 2: Score From iBench

We show the five runs of both cases in Table 2. For each run, the machine is clean boot for each run. The scores are quite close and show no difference on the system running with and without Argus.

We also perform the evaluation on benchmark of Chromium, while recoding the time usage. The telemetry from Chromium project is used to measure the CPU overhead of Argus. In the experiment, it uses the chromium to launch webpage and scroll over. The result is shown in Table 3.

	bare run	tracing run
real	27.7s	28.0s
user	28.3s	28.3s
sys	5.0s	5.7s

Table 3: Chromium benchmark: telemetry

As shown in Table 3, the overhead for real, user and sys are 1%, 0% and 14% respectively.

7. Related Work

While there is currently no system that can help users debug performance issues in closed-source applications on proprietary MacOS, several active research topics are closely related.

Event tracing. Magpie [5] is perhaps the closest to our work. It monitors server applications in Windows with the goal to model the normal behaviors of a server application in response to a workload. This model further helps detecting anomalies statistically. In contrast, Argus’s goal is to identify the root causes of performance issues. Its graphs are not request graphs, but rather graphs that may contain many requests. In addition, it logs normal and abnormal executions in the same event trace. In addition, Magpie requires a manual-written event schema for all involved applications to capture precise request graphs, whereas Argus has a simple, application-agnostic schema for system-wide tracing and enables users to provide more application-specific knowledge on demand.

Panappticon [23] monitors a mobile system and uses the trace to characterize the user transactions of mobile apps. Although it aims to track system-wide events and correlate them without developer input, it supports only two models of communication: work queue and thread pooling.

AppInsight [16] instruments application to identify the critical execution path in a user transaction. It supports the event callback pattern, and does not trace across process or app boundaries.

XTrace, Pinpoint and etc [11, 6, 7] trace the path of a request through a system using a unique identifier attached to each request and stitch traces together with the identifier. Argus does not assume the presence of a unified identifier in closed-source, third-party applications, frameworks, and libraries.

Aguilela [3] uses timing analysis to correlate messages to recover their input-output relations while treating the application as a black box.

Performance anomaly detection. Several systems detect performance anomalies automatically. [12, 22] leverage the user logs and call stacks to identify the performance anomaly.

[8, 18, 20, 9] apply the machine learning method to identify the unusual event sequence as an anomaly. [21] generates the wait and waken graph from sampled call stacks to study a case of performance anomaly.

These systems are orthogonal to Argus as Argus’s goal is to diagnose an already-detected performance anomaly. These systems can help Argus by detecting more accurately when a performance issue arises.

8. Conclusion

Our key insight in this paper is that causal tracing is inherently imprecise. We have reported patterns we observed that pose big precision challenges to causal tracing, and built Argus, a practical system for effectively debugging performance issues in MacOS applications despite the imprecision of causal tracing. To do so, it lets a user provide domain knowledge interactively on demand. Our results show that Argus effectively helped us locate all root causes of the issues, including a bug in Chromium, and incurred only 1% CPU overhead in its system-wide tracing.

References

- [1] 2008.
- [2] 2012.
- [3] Marcos K Aguilela, Jeffrey C Mogul, Janet L Wiener, Patrick Reynolds, and Athicha Muthitacharoen. Performance debugging for distributed systems of black boxes. In *ACM SIGOPS Operating Systems Review*, volume 37, pages 74–89. ACM, 2003.
- [4] Mona Attariyan, Michael Chow, and Jason Flinn. X-ray: Automating root-cause diagnosis of performance anomalies in production software. In *Presented as part of the 10th USENIX Symposium on Operating Systems Design and Implementation OSDI 12*, pages 307–320, 2012.
- [5] Paul Barham, Austin Donnelly, Rebecca Isaacs, and Richard Mortier. Using magpie for request extraction and workload modelling. In *OSDI*, volume 4, pages 18–18, 2004.
- [6] Mike Y Chen, Emre Kiciman, Eugene Fratkin, Armando Fox, and Eric Brewer. Pinpoint: Problem determination in large, dynamic internet services. In *Proceedings International Conference on Dependable Systems and Networks*, pages 595–604. IEEE, 2002.
- [7] Michael Chow, David Meisner, Jason Flinn, Daniel Peek, and Thomas F Wenisch. The mystery machine: End-to-end performance analysis of large-scale internet services. In *11th USENIX Symposium on Operating Systems Design and Implementation OSDI 14*, pages 217–231, 2014.
- [8] Ira Cohen, Jeffrey S Chase, Moises Goldszmidt, Terence Kelly, and Julie Symons. Correlating instrumentation data to system states: A building block for automated diagnosis and control. In *OSDI*, volume 4, pages 16–16, 2004.
- [9] Min Du, Feifei Li, Guineng Zheng, and Vivek Srikumar. Deeplog: Anomaly detection and diagnosis from system logs through deep learning. In *Proceedings of the 2017 ACM SIGSAC Conference on Computer and Communications Security*, pages 1285–1298. ACM, 2017.
- [10] Eun, 2018.
- [11] Rodrigo Fonseca, George Porter, Randy H Katz, Scott Shenker, and Ion Stoica. X-trace: A pervasive network tracing framework. In *Proceedings of the 4th USENIX conference on Networked systems design & implementation*, pages 20–20. USENIX Association, 2007.
- [12] Shi Han, Yingnong Dang, Song Ge, Dongmei Zhang, and Tao Xie. Performance debugging in the large via mining millions of stack traces. In *Proceedings of the 34th International Conference on Software Engineering*, pages 145–155. IEEE Press, 2012.
- [13] Tyler Harter, Chris Dragg, Michael Vaughn, Andrea C Arpaci-Dusseau, and Remzi H Arpaci-Dusseau. A file is not a file: understanding the i/o behavior of apple desktop applications. *ACM Transactions on Computer Systems (TOCS)*, 30:10, 2012.
- [14] Galen Hunt and Doug Brubacher. Detours: Binary interception of win 32 functions. In *3rd usenix windows nt symposium*, 1999.

- [15] Karthik Nagaraj, Charles Killian, and Jennifer Neville. Structured comparative analysis of systems logs to diagnose performance problems. In *Proceedings of the 9th USENIX conference on Networked Systems Design and Implementation*, pages 26–26. USENIX Association, 2012.
- [16] Lenin Ravindranath, Jitendra Padhye, Sharad Agarwal, Ratul Mahajan, Ian Obermiller, and Shahin Shayandeh. Appinsight: mobile app performance monitoring in the wild. In *Presented as part of the 10th USENIX Symposium on Operating Systems Design and Implementation OSDI 12*, pages 107–120, 2012.
- [17] Patrick Reynolds, Charles Edwin Killian, Janet L Wiener, Jeffrey C Mogul, Mehul A Shah, and Amin Vahdat. Pip: Detecting the unexpected in distributed systems. In *NSDI*, volume 6, pages 9–9, 2006.
- [18] Ali G Saidi, Nathan L Binkert, Steven K Reinhardt, and Trevor Mudge. Full-system critical path analysis. In *ISPASS 2008-IEEE International Symposium on Performance Analysis of Systems and software*, pages 63–74. IEEE, 2008.
- [19] Weiwei Xiong, Soyeon Park, Jiaqi Zhang, Yuanyuan Zhou, and Zhiqiang Ma. Ad hoc synchronization considered harmful. In *OSDI*, volume 10, pages 163–176, 2010.
- [20] Wei Xu, Ling Huang, Armando Fox, David Patterson, and Michael I Jordan. Detecting large-scale system problems by mining console logs. In *Proceedings of the ACM SIGOPS 22nd symposium on Operating systems principles*, pages 117–132. ACM, 2009.
- [21] Xiao Yu, Shi Han, Dongmei Zhang, and Tao Xie. Comprehending performance from real-world execution traces: A device-driver case. In *ACM SIGPLAN Notices*, volume 49, pages 193–206. ACM, 2014.
- [22] Ding Yuan, Soyeon Park, Peng Huang, Yang Liu, Michael M Lee, Xiaoming Tang, Yuanyuan Zhou, and Stefan Savage. Be conservative: enhancing failure diagnosis with proactive logging. In *Presented as part of the 10th USENIX Symposium on Operating Systems Design and Implementation OSDI 12*, pages 293–306, 2012.
- [23] Lide Zhang, David R Bild, Robert P Dick, Z Morley Mao, and Peter Dinda. Panappticon: event-based tracing to measure mobile application and platform performance. In *2013 International Conference on Hardware/Software Codesign and System Synthesis (CODES+ ISSS)*, pages 1–10. IEEE, 2013.
- [24] Xu Zhao, Kirk Rodrigues, Yu Luo, Ding Yuan, and Michael Stumm. Non-intrusive performance profiling for entire software stacks based on the flow reconstruction principle. In *12th USENIX Symposium on Operating Systems Design and Implementation OSDI 16*, pages 603–618, 2016.