

Argus : Debugging Performance Issues in Modern Applications with Interactive Causal Tracing

Abstract

Prior systems use causal tracing, a powerful technique that traces low-level events and builds request graphs, to diagnose performance issues. However, they all assume that accurate causality can be inferred from low-level tracing with supported communication patterns or obtained from developer’s schema upfront for all involved components. Unfortunately, based on our study and experience of building a causal tracing system on macOS, we find it is difficult, if not impossible, to get accurate request graphs. We present Argus, a practical system for debugging performance issues in modern desktop applications despite the inaccuracy of causal tracing. Argus lets a user inspect current diagnostics and provide domain knowledge on demand to counter the inherent inaccuracy of causal tracing. We implemented Argus in macOS and evaluated it on 11 real-world, open spinning-cursor issues in popular applications. The root causes of these issues were largely previously unknown. Our results show that Argus effectively helps us locate all root causes of the issues and incurs only 7% CPU overhead in its system-wide tracing.

1 Introduction

Today’s web and desktop applications are predominantly parallel or distributed, making performance issues in them extremely difficult to diagnose because the handling of an external request is often spread across many threads, processes, and asynchronous contexts instead of in one sequential execution segment [21]. To manually reconstruct a graph of execution segments for debugging, developers have to sift through a massive amount of log entries and potentially code of related application components [14, 24, 29, 31, 33]. More often than not, developers give up and resort to guessing the root cause, producing “fixes” that sometimes make the matter worse. For instance, a bug in the Chrome browser engine causes a spinning cursor in macOS when a user switches the input method [8], was first reported in 2012. Developers attempted to add timeouts to work around the issue. Unfortunately, the bug has remained open for seven years and the timeouts obscured diagnosis further.

Prior work proposed what we call *Causal tracing*, a powerful technique to construct request graphs automatically [13, 19, 25, 26, 32]. It does so by inferring (1) the beginning and ending boundaries of the execution segments (vertices in the graph) involved in handling a request; and (2) the causality between the segments (edges)—how a segment causes others to do additional handling of the request. Compared to debuggers such as *spindump* that capture only the current

system state, causal tracing is effective at aiding developers to understand complex causal behaviors and pinpoint the root causes for real-world performance issues.

Prior causal tracing systems all assume certain programming idioms to automate inference. For instance, if a segment sends a message, signals a condition variable, or posts a task to a work queue, it wakes up additional execution segment. Prior systems assume that wake-ups reflect causality. Similarly, they assume that the execution segment, from the beginning of a callback invocation to the end, is entirely for handling related work in a request. Unfortunately, based on our study and experience of building a causal tracing system for macOS, we find that modern applications violate these assumptions. Hence, the request graphs computed by causal tracing are inaccurate in several ways.

First, an inferred segment can be larger than the actual event handling segment due to batch processing. Specifically, for performance, an application or its underlying frameworks may bundle together work on behalf of multiple requests with no clear distinguishing boundaries. For instance, WindowServer in macOS sends a reply for a previous request and receives a message for the current request using one system call *mach_msg_overwrite_trap*, presumably to reduce user-kernel crossings. Second, the graphs can miss numerous causal edges. For instance, consider data dependencies in which the code sets a flag (e.g., “*need_display = 1*” in macOS animation rendering) and later queries the flag to process a request further. This pattern is broader than ad hoc synchronization [28] because data dependency occurs even within a single thread (such as the buffer holding the reply in the preceding WindowServer example). Although the number of these flags may be small, they often express critical causality, and not tracing them would lead to many missing edges in the request graph. However, without knowing where the flags reside in memory, a tool would have to trace all memory operations, incurring prohibitive overhead and adding many superfluous edges to the request graph. Third, inferred edges can be superfluous because wake-ups do not necessarily reflect causality. Consider an *unlock()* operation waking up an thread waiting in *lock()*. This wake-up can just be happenstance and the developer’s intent is mutual exclusion. However, the two operations can also enforce a causal order.

We believe that, without fully understanding of application semantics, request graphs computed by causal tracing are *inherently* inaccurate and both over- and under-approximate reality. Although developer annotations can help improve accuracy [19, 26], modern applications use

more and more third-party libraries, whose source code is not available. Given the frequent use of custom synchronizations, work queues, and data flags in modern applications, it is hopeless to count on manual annotations to ensure accurate capture of request graphs.

In this work, we present Argus, a dramatically different approach in the design space of causal tracing. It is designed for tech-savvy users who are interested in compiling useful bug reports for their daily use applications, whose source codes are often not available. As opposed to full manual schema upfront for all involved applications and daemons [12, 19, 26], Argus calculates an event graph for a duration of the system execution with both true causal edges and weak ones, and enables users to provide necessary schematic information on demand in diagnosis. Specifically, it keeps humans in the loop, as a debugger should rightly do. Argus queries users a judiciously few times to (1) resolve a few inaccurate edges that represent false dependencies and (2) identify potential dependency due to data flags.

We implement Argus in macOS, which is closed-source on its frameworks and many applications. This closed-source environment therefore provides a true test of Argus. We address multiple nuances of macOS that complicate causal tracing, and build a system-wide, low-overhead, always-on tracer. Argus enables users to optionally increase the granularity of tracing (e.g., logging call stacks and instruction streams) by integrating with existing debuggers such as *lldb*.

We evaluate Argus on 11 real-world, open spinning-cursor issues in widely used applications such as Chromium browser engine and macOS System Preferences, Installer, and Notes. The root causes of all 11 issues were previously unknown to us and, to a large extent, the public. Our results show that Argus is effective: it helps us non-developers of the applications find all root causes of the issues, including the Chromium issue that remained open for seven years. Argus mostly needs only less than 3 user queries per issue but they are crucial in aiding diagnosis. Argus is also lightweight: its systems-wide tracing incurs only 7% CPU overhead overall.

We make the following contributions:

1. We demonstrate conceptual realization that causal tracing is inherently inaccurate, and introduce interactive approach in the design space of causal tracing.
2. We build Argus, performing scalable system-wide causal tracing with little overhead, and design its diagnosis algorithm with comparison to normal cases.
3. We use it to diagnose real-world spinning cursors and find root causes for performance issues that have remained open for years.

This paper is organized as follows. In Section 2, we introduce the causal tracing and prior works. In Section 3, we present an overview of using Argus and a Chromium

example. In Section 4, we report inherently inaccurate patterns observed in macOS, and Section 5 describes our tracing implementation and tools for user interaction. Section 6 demonstrates the methodology and results of case studies, and the performance evaluation. We summarize related work in Section 7, and end with conclusion in Section 8.

2 Background

In this section, we illustrate causal tracing and prior work on it. Causal tracing collects events standing for instructions executed in CPU and generates a graphical representation with traced events as vertices. Two events always following a sequential constraint reflects causality, which is represented as an edge. The graph helps users understand the complex causal behaviors across thread/process boundaries and attribute bugs to their root causes. Prior works have different definitions of vertices, edges, and root causes based on what events are collected.

AppInsight instruments all the upcalls from the framework to the application. It traces user input, display update, the begin and end of procedural call, the invocation of callback function, exception and blocking events in threads. Each event is reflected as a vertex in the request graph. Therefore, the request graph connects (1) user input event to (2) the beginning of event handler, which in turn connects to (3) the beginning of callback in background threads. The vertices like (2) and (3) will connect to (4) the end of the procedural call or lead to (5) exception. Besides, they will also connect to (6) blocking for signal, if the execution requires synchronization. The goal of AppInsight is to help developers understand the performance bottlenecks with critical paths or exception paths. It defines the root cause as the state of a function execution, long blocking or exception in the application.

To be unobtrusive, Panappticon instruments the system to collect low-level and fine grained events from libraries and kernel, including user input, display update, asynchronous call and callback, inter-process communication, synchronization mechanism, and resource accounting. Every event is a vertex. Panappticon connects continuous vertices which stem from atomic work in a thread, e.g., a worker thread processing one task from a task queue, into an execution interval. Two execution intervals are connected if the earlier interval triggers the latter one. For example, a user input triggering an enqueue message in the same thread reflects as an execution interval, where two vertices are connected with a temporal ordering edge. In another thread, dequeuing the message and submitting an asynchronous task generate another execution interval. The two intervals are connected with a causal edge. With the resources analysis in every user transaction, from user input to display update, Panappticon speculates and manually inspect root causes from design flaws, harmful interaction, to underpowered hardware.

3 Overview

In this section, we first describe Argus event graphs (§3.1). Then we explain how Argus diagnoses performance anomaly with a concrete example (§3.2).

3.1 Event Graph Basics

The event graph is a generalized control-flow graph which includes inter-thread and inter-process dependencies. To construct event graphs, Argus collects three categories of events in its systems-wide event logs. The first category of events are boundary events that mark the beginning and ending of execution segments. Argus handles common callbacks, such as *dispatch_client_callout* and *CFRunLoopDoBlocks*, and mark their entry and return as boundaries. Every execution segment corresponds to a vertex in the event graph. The second category contains semantic events, including system calls, call stacks when certain operations such as *mach_msg* are running, and user actions such as key presses. These events are stored as contents in the vertex and primarily for providing information to user during diagnosis. The third category of events are communication events for forming edges in the graph. For instance, an operation that installs a callback is connected to the invocation of the callback.

A unique design in Argus is tracing general wake-up and wait operations inside the kernel to ensure coverage across diverse user-level custom synchronization primitive because their implementations almost always use kernel wake-up and wait. Argus includes the wake-up in the category of communication event, and classifies the wait into the category of boundary event. This approach necessarily includes spurious edges in the graph, including those due to mutual exclusion and context switch by interrupts; Argus handles them by querying the user when it encounters a vertex with multiple incoming causal edges during diagnosis (see §3.2). We also observed that a waiting kernel thread is frequently woken up to perform tasks such as timer firing signal and scheduler maintenance; Argus recognizes them and culls them out from the graph automatically.

Compared to tools such as *spindump* that capture only the current system state, event graphs capture the causal path of events, enabling users to trace across threads and processes to events happened in the past (hence cannot be captured by *spindump*) that explain present anomalies. Therefore Argus can report root causes such as dead locks due to design flaws.

3.2 Argus Work Flow

In this section, we describe the steps a tech-savvy user takes to investigate a performance anomaly with Argus. Figure 1 shows Argus’s work flow with an example of a user investigating a performance problem in Chromium. The system wide tracing tool, which collects data from Argus instrumented library and kernel, generates logs. They are transformed into an *event graph* in Argus’s graph construction

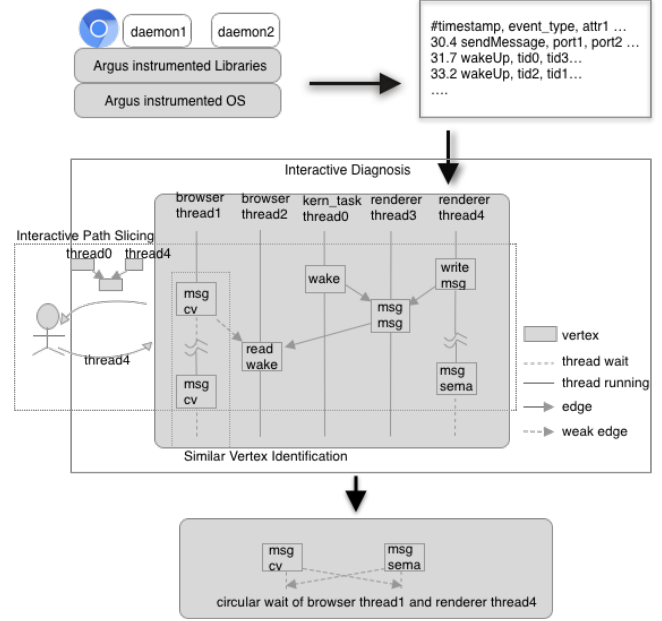


Figure 1. Argus Work Flow

component. Diagnosis and inference are performed within this graph, in a semi-automated fashion. As shown in Figure 1, the Argus interactive debugger asks the user to choose one edge in a subgraph. With this type of interaction, Argus runs its diagnosis algorithm and reports the root cause vertices. Next, we describe how Argus assists the user to diagnose a performance issue.

3.2.1 Initiate the diagnosis

Compared to prior works which try to identify performance issue with critical paths, Argus takes the reported issue by system or users as input. It leverages the semantic events in the graph to locate the vertex representing the reported issue and initiate the diagnosis.

Consider a common performance bug on macOS, the *spinning cursor*, which indicates the current application’s main thread has not processed any UI events for over two seconds. *WindowServer* throws *spinning cursor* with the collaboration of the application’s *NSEvent* thread. The *NSEvent* thread fetches *CoreGraphics* events from *WindowServer*, converts and creates *NSApp* events for the main thread, and arms a timer to monitor the event’s wait time. If the main thread fails to process an event before the timer fires, the *NSEvent* thread notifies *WindowServer* via “*CGSConnectionSetSpinning*”, and *WindowServer* draws a *spinning cursor*.

To initialize debugging a *spinning cursor*, Argus first constructs an event graph from the system-wide event log. It then checks the semantic events from the *NSEvent* thread, identifies the time when the performance anomaly happens, the invocation of “*CGSConnectionSetSpinning*”, and finds

spinning vertex in the main thread which contains the ongoing events concurrent.

3.2.2 Diagnosis with Graph

Algorithm 1 Diagnosis algorithm.

Input: g - EventGraph; spinning_vertex- the vertex in the UI thread when the spinning cursor occurs

Output: root_cause_vertices- vertices for user to inspect root cause

```

1: function DIAGNOSE( $g$ , spinning_vertex)
2:   switch spinning_vertex.block_type do
3:     case LongRunning
4:       slice  $\leftarrow$  InteractiveSlice(spinning_vertex)
5:       return vertex contains UI event
6:     case RepeatedYield
7:       if DataFlagEvent  $\notin$  {event types in spinning_vertex} then
8:         Require users to annotate data flag
9:         abort()
10:      end if
11:      /* Fall through */
12:     case LongWait
13:       similar_vertex  $\leftarrow$  vertex has similar event sequence to spinning_vertex
14:       normal_path  $\leftarrow$  InteractiveSlice(similar_vertex)
15:       for each  $t \in$  {threads in normal_path} do
16:          $vertex_t \leftarrow$  vertex in  $t$  before spinning_vertex spins
17:         if  $vertex_t \in$  {LongRunning, RepeatedYield, LongWait} then
18:           root_cause_vertices.append( $vertex_t$ )
19:           root_cause_vertices.append(DIAGNOSE( $g$ ,  $vertex_t$ ))
20:         end if
21:         /* if  $t$  is normal running, diagnose the next thread */
22:       end for
23:     end switch
24:     return root_cause_vertices
25: end function

26: function INTERACTIVESLICE( $g$ , vertex)
27:   loop
28:     path_slice.append(vertex)
29:     if vertex has 1 incoming causal edge then
30:       vertex  $\leftarrow$  predecessor vertex
31:     else if vertex has multiple incoming edges then
32:       vertex  $\leftarrow$  ask user to pick from predecessors
33:     else if vertex had weak edge then
34:       vertex  $\leftarrow$  ask user to pick or stop
35:     else
36:       /* The first vertex of current thread */
37:       return path_slice
38:     end if
39:     if vertex is invalid then
40:       /* user chooses to stop traversal with invalid input */
41:       return path_slice
42:     end if
43:   end loop
44: end function

```

Given the event graph and the spinning vertex, Argus runs Algorithm 1 to pinpoint the root cause. Specifically, upon examining what the main thread is actually doing, there are three potential cases.

- **LongRunning** (lines 3 - 5). The main thread is busy performing lengthy CPU operations. This case is the simplest, and Argus traverses the event graph backwards to find a slice originating from the offending UI event to the long running CPU operations. This slice is particularly useful for further diagnosing the bug. As shown in Function *InteractiveSlice* in line 26, Argus may encounter vertices with multiple incoming edges or weak edges that do not reflect causality when traversing the graph. It queries the user to resolve them.
- **RepeatedYield** (lines 6 - 11). The main thread is in a yield loop, which is highly indicative it is waiting on a data flag (e.g., “while(!done) thread_switch();”). If Argus cannot find any record of data flags in the spinning vertex, it terminates debugging by prompting the user to identify data flags and re-trace the application. Here we assume that the performance issue reproduces with a reasonable probability because, fortunately, a one-off issue that never reproduces is not as annoying as one that occurs frequently. If Argus finds the data flag the spinning vertex is waiting for, it falls through to the next case.
- **LongWait** (lines 12 - 22). The main thread is in a lengthy blocking wait and the wake-up has been missing. Argus handles this case by finding a normal scenario where the wake-up indeed arrives, and then figures out which wake-up edge is missing in the spinning scenario along the expected wake-up path. Specifically, Argus first finds a similar vertex to the spinning one based on the event sequences in each vertex(3.3). It then traverses backwards from the similar vertex to find the normal wake-up path. For each thread in the wake-up path, it examines the vertex in the thread right before the spinning vertex waits. If this vertex is also abnormal, Argus appends it to the path of root cause vertices, and applies Function *Diagnose* recursively to diagnose “the culprit of the culprit”. For each such vertex, it queries the user to determine whether to proceed or stop, because based on our experience a few iterations are adequate to manifest the root cause.

Our experience suggests that the first case is the most common, but the second and third represent more severe bugs. Long-running CPU operations tend to be more straightforward to diagnose with existing tools such as *spindump* except they do not connect CPU operations back to daemons and UI events. Repeated yielding or long waiting cases involve multiple threads and processes, and are extremely hard to understand and fix even for the application’s original developers. Therefore, issues remain unaddressed for years and significantly impact the user experience. Algorithm 1 is semi-automated but can integrate user input to leverage hypotheses or expert knowledge as to why a hang may occur.

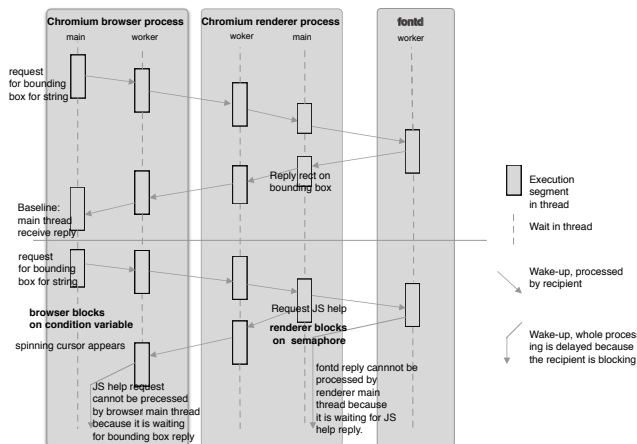


Figure 2. Chromium Spinning Cursor Example

Our results show that user inputs, albeit few, are crucial in this process (§6.3).

3.2.3 Chromium Spinning Cursor Example

One of the authors experienced first-hand performance issue in Chromium, an open-source browser engine that powers Google Chrome and, starting recently, Microsoft Edge [7]. She tried to type in the Chromium search box a non-English word with a default Input Method Editor shipped with macOS. The browser appeared frozen and the spinning cursor showed for a few seconds. Afterwards everything went back to normal. This issue is reproducible and always ruins her experience, but is quite challenging to diagnose because two applications Chromium and IME and many daemons ran and exchanged messages. It was reported by other users for other non-English input methods, too.

To diagnose this issue with Argus, the author follows the steps in Figure 1. She started system-wide tracing, and then reproduced the spinning cursor with a non-English search string while the page was loading. After the very first few characters which the browser handles normally, the remaining characters triggered a spinning cursor. The entire session took roughly five minutes. She then ran Argus to construct the event graph. The graph was highly complex, with 2,749,628 vertices and 3,606,657 edges, almost fully connected. It spanned across 17 applications; 109 daemons including *fontd*, *mdworker*, *nsurlsessiond* and helper tools by applications; 126 processes; 679 threads, and 829,287 messages. Given the scale of the graph and the diverse communication patterns, it would be extremely challenging for prior automated causal tracing tools [10, 11, 16, 32] because they handle a limited set of patterns. Tools that require manual schema [12, 26], would be prohibitive because developers would have to provide schema for all involved applications and daemons.

Next she ran Argus to find the spinning vertex in the main thread of the browser process. Argus returned a *LongWait* event, a *psynch_cv_wait()* with a timeout of 1.5 seconds, and identified a similar vertex in normal scenario where the *wait* was signaled quickly.

Argus then found the normal wake-up path with interactive path slicing, which connects five threads. The *browser* main thread was signaled by a *browser* worker thread, which received IPC from a worker thread of *renderer* where the rendering view and WebKit code run. This worker thread is woken up by the *renderer* main thread, which in turn woken by *fontd*, the font service daemon, as shown in the upper side of Figure 2. Argus further compared the normal case with the spinning case thread by thread as shown in Figure 2, and returned the *LongWait* event on semaphore in the *renderer* main thread, the culprit that delayed waking up the *browser* main thread over four seconds. What caused the wait in the *renderer* main thread though? She thus continued diagnosis and recursively applied Argus to the wait in *renderer* (corresponding normal path is not shown in the figure), and it turned out that *renderer* was actually waiting for the *browser* main thread – a circular wait formed.

Inspect of the reported vertices reveals that the *browser* was waiting for the *renderer* to return the string bounding box for a string, and the *renderer* was waiting for the *browser* to help render JavaScript. This circular wait was broken by a timeout in the browser main thread (the *wait* on *psynch_cv_wait()* timeouts in 1,500 ms). While the system was able to make progress, the next key press caused the spinning cursor to display for another 1,500 ms. The timeout essentially converted a deadlock into a livelock. To verify this diagnosis, we shortened the timeout in the *psynch_cv_wait()* called in the Chromium browser to 150 ms, which proportionally reduced how often this spinning cursor occurs.

3.3 Similar Vertex Indetification

Argus leverages a normal scenario to discover missing wake-up edges in *LongWait* cases. The normal scenario is identified from the similar vertex which shares the same high level semantics as the spinning vertex, but exposes different execution results.

Argus identifies the similar vertices first by inferring its high level semantics with semantic events, and its programming purpose with communication events which form edges in the graph. The semantic events include system calls, call stacks, and user inputs. Their sequential order and runtime unrelated attributes are treated as hallmarks. For example, in call stacks, Argus compares vertices with symbol names instead of memory addresses. The communication events include messages, dispatch queue operations, runloop operations, *data_flag_read* and *data_flag_write*. The process names of peer vertices are checked, as they reflect the developers' intent, for example, to request service from a particular

```

1 //worker thread in fontd:      1 //main thread in fontd:
2 //enqueue a block             2 //dequeue blocks
3 block = dispatch_mig_sevice;   3 block = dequeue();
4 dispatch_async(block);         4 dispatch_execute(block);

1 //implementation of dipatch_mig_server
2 dispatch_mig_server()
3 for(;;) //batch processing
4     mach_msg(send_reply, rcv_request)
5     call_back(rcv_request)
6     set_reply(send_reply)

```

Figure 3. Dispatch message batching

helper tool. More over, a user can also instruct Argus to inspect preceeding vertices to improve identification accuracy. Then Argus discriminates the execution results to filter out spinning scenarios. with the time cost and wait result in the vertex by default. If multiple similar vertices are identified, Argus usually chooses the most recent one heuristically.

3.4 Limitations

Argus is designed to support interactive debugging of performance issues. It requires a user to reproduce a performance issue, in the following situations: (1) Argus need capture more fine-grained event traces like accesses to data flags to mitigate under-connection of event graph; (2) The root cause vertices reported by Argus lack high-level semantics, so Argus need integrate *llbd* to gather more information like call stacks. Fortunately, a performance issue that never reproduces is probably not as annoying as one that occurs now and then. On the other hand, the blocking issue always showing up and never signaled normally are out of the debugging range of Argus, because Argus’s diagnosis algorithm never finds its comparable normal causal path.

We implemented Argus in the closed-source macOS which presents a harsh test, but have not ported Argus to other operating systems yet. It is possible that the techniques do not generalize to other operating systems. However, modern operating systems share many similarities, and inspire each others’ designs, so we are hopeful that the ideas are generally applicable.

4 Handling Inaccuracies in Causal Tracing

4.1 Inherent Inaccuracies

As explained in §2, causal tracing builds a graph to connect execution segments on behalf of a request that spread across separate threads and processes. Based on our experience of building a causal tracing system on commercial closed-source macOS, we believe such graphs are inherently inaccurate and contain both *over-connections* which do not really map to causality, and *under-connections*, missing edges between two vertices with one causally influencing the other.

```

1 //inside a single thread
2 CGXPostReplyMessage(msg) {
3     // send _gOutMsg if it hasn't been sent
4     push_out_message(_gOutMsg)
5     _gOutMsg = msg
6     _gOutMessagePending = 1
7 }
8 CGXRunOneServicePass() {
9     if (_gOutMessagePending)
10        mach_msg_overwrite(MSG_SEND | MSG_RECV, _gOutMsg)
11    else
12        mach_msg(MSG_RECV)
13    ... // process received message
14 }

```

Figure 4. Batching in event processing

```

1 //Worker thread:      1 //Main thread:
2 //needs to update UI: 2 //traverse all CA objects
3 obj->need_display = 1  3 if(obj->need_display == 1)
                        4     render(obj)

```

Figure 5. CoreAnimation shared flag

```

1 //NSEvent thread:
2 CGEventCreateNextEvent() {
3     if (sCGEventIsMainThreadSpinning == 0x0)
4         if (sCGEventIsDispatchToMainThread == 0x1)
5             CFRRunLoopTimerCreateWithHandler{
6                 if (sCGEventIsDispatchToMainThread == 0x1)
7                     sCGEventIsMainThreadSpinning = 0x1
8                     CGSCConnectionSetSpinning(0x1);
9             }
10 }

1 //Main thread:
2 {
3     ... //pull events from event queue
4     Convert1CGEvent(0x1);
5     if (sCGEventIsMainThreadSpinning == 0x1){
6         CGSCConnectionSetSpinning(0x0);
7         sCGEventIsMainThreadSpinning = 0x0;
8         sCGEventIsDispatchedToMainThread = 0x0;
9     }
10 }

```

Figure 6. Spinning Cursor Shared Flags

In this section, we present several inherently inaccurate patterns we observed and their examples.

4.1.1 Over Connections

Over connections usually occur when (1) intra-thread boundaries are missing due to unknown batch processing programming paradigms or (2) superfluous wake-ups that do not always imply causality.

Dispatch message batching While traditional causal tracing assumes the entire execution of a callback function is on behalf of one request, we found some daemons implement

their service loop inside the callback function and create false dependencies. In the code snippet in Figure 3 from the *fontd* daemon, function *dispatch_client_callout* is installed as a callback to a work from dispatch queue. It subsequently calls *dispatch_mig_server()* which runs the typical server loop and handles messages from different apps. Any application or daemon can implement its own server loop this way, which makes it fundamentally difficult to automatically infer event handling boundaries.

Batching in event processing Message activities inside a system call are assumed to be related traditionally. However, to presumably save on kernel boundary crossings, WindowServer MacOS system daemon uses a single system call to receive data and send data for an unrelated event from different processes in its event loop in Figure 4. This batch processing artificially makes many events appear dependent.

Mutual exclusion In a typical implementation of mutual exclusion, a thread’s unlock operation wakes up a thread waiting in lock. Such a wake-up may be, but is not always, intended as causality. However, without knowing the developer intent, any wake-up is typically treated as causality.

4.1.2 Under Connections

We observe that under connections mostly result from missing data dependencies. This pattern is more general than shared-memory flags in ad hoc synchronization [28] because it occurs even within a single thread.

Data dependency in event processing The code for Batching in event processing above also illustrates a causal linkage caused by data dependency in one thread. WindowServer saves the reply message in variable *_gOutMsg* inside function *CGXPostReplyMessage*. When it calls *CGXRunOneServicePass*, it sends out *_gOutMsg* if there is any pending message.

CoreAnimation shared flags As shown in the code snippet Figure 5, worker thread can set a field *need_display* inside a *CoreAnimation* object whenever the object needs to be repainted. The main thread iterates over all animation objects and reads this flag, rendering any such object. This shared-memory communication creates a dependency between the main thread and the worker so accesses to these field flags need to be tracked.

Spinning cursor shared flag As shown in Figure 6, whenever the system determines that the main thread has hung for a certain period, and the spinning beach ball should be displayed, a shared-memory flag is set. Access to the flag is controlled via a lock, i.e. the lock is used for mutual exclusion, and does not imply a happens before relationship.

4.2 Handling Inaccuracies

In this section, we discuss how Argus leverage heuristics to mitigate over-connections and under-connections.

4.2.1 Mitigating Over-Connections

From a high-level, Argus deals with over-connections by heuristically splitting an execution segment that appears mixing handling of multiple requests.

Specifically, Argus splits based on three criteria. First, Argus recognizes a small set of well-known batch processing patterns such as *dispatch_mig_server()* in §4 and splits the batch into individual items. Second, when a wait operation such as *recv()* blocks, Argus splits the segment at the entry of the blocking wait. The rationale is that blocking wait is typically done at the last during one step of event processing. Third, if a segment communicates to too many peering processes, Argus splits the segment when the set of peers differs. Specifically, for each message, Argus maintains a set of two peers including (1) the direct sender or receiver of the message and (2) the beneficiary of the message (macOS allows a process to send or receive messages on behalf of a third process). Argus splits when two consecutive message operations have non-overlapping peer sets.

4.2.2 Mitigating Under-Connections

Causal tracing infers boundaries for the background thread with blocking wait. It can break an "atomic" task in the worker thread, e.g. a thread pauses for IO resource. Argus adds weak causal edges to mitigate under-connection. When a weak edge is encountered during diagnosis, it queries users to decide whether to follow the weak edge or stop (§3).

The other primary under-connections are due to data dependencies. Currently Argus queries the user to identify the data flags. It is conceivable to leverage memory protection techniques to infer them automatically, as demonstrated in previous record-replay work [18, 23]. It is out of the scope of this paper and we leave it for future work. Currently, to discover a data flag, the user re-runs the application with Argus to collect instruction traces of the concurrent events in both the normal and spinning cases and detects where the control flow diverges. Argus exposes register values for the basic blocks before the divergence and uncovers the address of the data flag. Once the user identifies a data flag, Argus traces it using either binary instrument, such as the *need_display* flag in *CoreAnimation* (§4), or with Argus’s watchpoint tool. Argus add a causal edge from a write to a data flag to its corresponding read.

5 Implementation

In this section, we discuss how Argus collects tracing events from both kernel and libraries.

5.1 Event Tracing

Current macOS supports a system-wide tracing infrastructure [1]. It by default stores events in memory and flushes them to screen or disk periodically. Argus extends this infrastructure with a file-backed ring buffer to support larger-scale

tests and avoid exhausting the disk. The file size is set to 2GB by default, approximately 19 million events (about 5 minutes of normal operations). The size can be adjust by users to accommodate the desired trace history.

Default tracing points in macOS provide too limited information to enact causal tracing. We therefore instrument both the kernel [2] (at the source level) and key libraries (at the binary level) to gather adequate data. We instrument the kernel with 1,193 lines of code, and binary-instrument the following libraries: *libsystem_kernel.dylib*, *libdispatch.dylib*, *libpthread.dylib*, *CoreFoundation*, *CoreGraphics*, *HIToolbox*, *AppKit* and *QuartzCore* in 57 different places.

5.2 Instrumentation

Most libraries and applications are closed-source in macOS. To hook their functions with techniques such as library preloading are not applicable, as libraries use two-level executable namespace [3]. Hence, we implemented a binary instrumentation mechanism that allows users to add tracing points inside a binary image.

Like Detour [22], we use static analysis to decide where the instrumentation performs and enact it at runtime. The user supplies a sequence of instructions for Argus to search the locations of interest in the image, and a trampoline function which overwrites the sequence of instructions and produces tracing data with API *kdebug_trace* from Apple. Argus generates shell code with the trampoline function to replace the victimized instructions. All of the trampoline functions are grouped into a new image with an initialization function that triggers the drop-in replacement. Argus configures the image in the following steps to finish the instrumentation: (1) it re-exports all symbols from the original image, so the original code can be called like a shared library; (2) it renames the original image and applies original name to the new one to ensure the modifications are properly loaded; (3) it invokes the initialization function externally with *dispatch_once* when the library loads.

5.3 Tracing Data Flags

As described in (§4), under-connection due to the missing data dependency requires users’ interaction. Users specify that reads and writes to a given variable should be considered data dependencies. Data flag tracing is possible through Argus’s binary rewriting, but we also provide a simple command line tool which uses watchpoint registers to record *data_flag_write* and *data_flag_read* events. The tool *bp_watch* takes as input the process ID, path to the relevant binary image, and the symbol name of the global variable. Here is a simple example of how a user asks Argus to trace *_gOutMsgPending*:

```
bp_watch Pid/of/WindowServer Path/to/CGs _gOutMsgPending
```

Argus loads the watchpoint handler into the address space of the target application by hooking it to *CoreFoundation*.

The handler invokes Apple’s API *kdebug_trace* to record the value of the data flag and its operation type (read or write).

5.4 Tracing Instructions and Calls

Users may need to gather more information, such as individual instructions and call stacks, to come up with and verify a binary patch. Argus integrates with *lldb scripts* to capture this information and add it to the corresponding vertices in the event graph. Our *lldb scripts* gather call stacks at relevant locations and parameters tainted by user inputs. To reduce the data collection overhead, only at each beginning of a function call does the script record a full call stack. While *lldb* steps into functions from apps and frameworks to record parameters tainted, it steps over and only records the return value of APIs from low-level libraries (i.e. those with the filename extension *.dylib*).

The combination of instruction-level tracing and occasional call-stacks offers more than enough detail to diagnose even the most arcane issues, and in our experience has been very helpful in multiple steps of an Argus diagnosis.

6 Evaluation

In this section, we describe the methodology of case study, the root causes diagnosed with Argus for performance issues in macOS, and the overhead of Argus tracing system.

6.1 Methodology

Prior work relies on the request graph per transaction to identify bottleneck and speculate possible causes. Panappticon for Android is one and closest to ours in design. Their traced events and causality are in a subset of ours. It is common to see multiple user inputs are connected into a Panappticon’s request graph, but the programming paradigm caused the over-connection is hard to reveal due to the closed-source Cocoa framework in macOS. We therefore carry out the inaccuracy study in §6.2 with a non-GUI toy program to demonstrate the inherent inaccuracy of causal tracing.

For real-world case study, we collect performance issues of popular applications, because they likely represent the bugs attractive to tech-savvy users. Among the 26 bugs from the github reports, we reproduced 3 of them successfully. Others are failed either due to the version capacity in ElCapitan or insufficient information in the bug report. 8 performance issues are collected from the daily use applications in the author’s laptop. As a result, we study 11 reproducible cases.

Now we describe how we measure the effects of heuristics which Argus uses to mitigate the graph inaccuracy, and manual efforts required in the diagnosis. We first enable tracing component and reproduce the performance issues in macOS. When Argus constructs event graph with the trace log, we measure the number of over-connection and under-connection mitigated by heuristics. After this step, a vertex in the graph can still have multiple incoming edges or weak


```

1 //insert debug info begin
2 ABAddressBook *AB = [ABAddressBook sharedAddressBook];
3 ABSearchElement *nameIsSmith =
4     [ABPerson searchElementForProperty:kABLastNameProperty...];
5 NSArray *ret = [AB recordsMatchingSearchElement:nameIsSmith];
6 //insert debug info end

```

Figure 7. Code snippet: access contacts

edge. We then run Argus diagnosis algorithm with a human in the loop, and count the times when multiple incoming edges or weak edges are encountered. Users can query the event graph for assistance or make decisions with domain knowledge. In the worst case that users make a wrong decision, before reaching the end of path slicing, Argus allows them to relocate the path to a particular vertex.

6.2 Inaccuracy study

In the section, we demonstrate the inaccuracy of request graph extracted with causal tracing and the necessary interactions in path slicing with study on a toy program as shown in Figure 7. It accesses the *Contacts* and search with last name *Smith*. We inject two tracing data in the source code at the beginning and end to mimic an end-to-end transaction.

The inherent inaccuracy of request graph is exposed with ambiguous edges. The asynchronous calls are always connected to their invocations in Panappticon. However, some of them are hard to decide. In the toy program, the causality between timer create and timer cancel introduces unrelated applications, e.g. *mds* and *trace* to the request graph. Removing this kind of edges on the other hand breaks the connections of event traces produced by line 2 and line3. Thus the request graph becomes incomplete in Panappticon. Argus mitigates the under-conneciton with the tracking of shared flags, *AB* and *nameIsSmith*.

6.3 Case Studies

In this section, we demonstrate how Argus helps to diagnose 11 spinning-cursor cases in popular applications. Table 1 describes these spinning-cursor cases. In Table 2, we compare Argus with Panappticon and list the portions of over- and under- connections mitigated with our hueristics. However, our filtered graph remains too inaccurate to automate diagnosis. The user interaction is still required but not overwhelming, as shown in Table 3. In most cases, up to 3 user queries suffice to find root cause path accurately. Although complex applications like MicosoftWord and Chromium require more queries, 13 and 22 respectively, many of them result from repeated patterns. They can be easily identified by users.

6.3.1 Long Wait and Repeated Yield

In this section, we discuss the cases where the spinning vertex is blocking on wait event or yielding loop, corresponding

Bug ID	Application	Bug description
1-Chromium	Chromium	Typing non-english in search box causes webpage freeze.
2-SystemPref	System Preferences	Disabling an online external monitor and rearranging windows causes System Preferences freeze.
3-SequelPro	Sequel Pro	Lost connection freezes the APP.
4-Installer	Installer	Moving cursor out of an authentication window causes freeze.
5-TeXStudio	TeXStudio	Modification on bib file with vim causes its main window hang.
6-TextEdit	TextEdit	Copying text over 30M causes freeze.
7-MSWord	Microsoft Words	Copying a document over 400 pages causes hang.
8-Notes	Notes	Launching Notes where stores a long note before causes freeze.
9-11 Others	SublimeText TextMate CotEditor	Copying or pasting in a file with large amount of context causes freeze.

Table 1. Bug Descriptions. We assign each bug in Column Bug ID to ease discussion

Bug ID	% over-connections	% under-connections	user provided data flag
1-Chromium	0.02	0.02	0
2-SystemPref	0.56	2.48	2
3-SequelPro	0.49	0.35	0
4-Installer	4.39	2.83	0
5-TeXStudio	2.43	0.58	0
6-Notes	2.97	11.53	0
7-TextEdit	7.97	0.72	0
8-MSWord	6.72	1.04	0
9-SIText	4.07	0.92	0
10-TextMate	2.15	2.18	0
11-CotEditor	4.81	5.32	0

Table 2. Inaccuracy handling compared to Panappticon

We calculate the over-connections mitigated with our splitting heuristics, and the under-connections where wait on primitive is treated as the end boundary, which Panappticon applies to background threads, but breaks an "atomic" task of a common callout. The percentages are calculated over the total edges in event graph.

to LongWait and RepeatedYield. These root causes are mostly can be verified by themselves, as they are manifested by comparing to normal scenarios.

2-SystemPref *System Preferences* provides a central location in macOS to customize system settings, e.g. additional monitors configuration. *DisableMonitor* [9] provides more functionality, enable/disable monitors online. We caught the spinning cursor while we disable an external monitor and rearrange windows in *Display* panel.

The log collected by Argus contains 2 cases: 1) a baseline scenario where the displays are rearranged with the enabled external monitor, and 2) a spinning scenario as we described above. The spinning vertex in the main thread is dominated by system calls, *mach_msg* and *thread_switch*, which falls

Bug ID	user interactions	size of path with		hueristic slicing over
		interaction	heuristics	interactive slicing
1-Chromium	13	32	303	9.47
2-SystemPref	1	2	30	15.00
3-SequelPro	2	5	264	52.80
4-Installer	2	6	36	6.00
5-TeXStudio	3	6	44	7.33
6-TextEdit	3	21	21	1.00
7-MSWord	22	67	136	2.03
8-Notes	2	10	42	4.20
9-SIText	1	3	3	1.00
10-TextMate	0	3	3	1.00
11-CotEditor	1	4	6	1.50

Table 3. Path slicing for "buggy" cases

we noticed that Argus’s query on incoming edges can be answered with the most recent one, thus we automated the path slicing in Argus with the heuristics. The length of the path is usually much longer due to excessive traverses to daemons and kernel task.

into the category of Repeated Yield. We discovered data flags, “_gCGWillReconfigureSeen” and “_gCGDidReconfigureSeen”, which signify the configuration status and break the thread-yield loop. Argus reveals that the main thread of *System Preferences*, in the baseline scenario, sets the flags after receiving specific datagrams from *WindowServer*. Conversely, the setting of “_gCGDidReconfigureSeen” is missing in the spinning case, and the main thread thus repeatedly sent messages to *WindowServer* for datagram.

In conclusion, we discovered that the bug is inherent in the design of the *CoreGraphics* framework, and would have to be fixed by Apple. We also verified this diagnosis by creating a dynamic binary patch to fix the deadlock. The patched library makes *DisableMonitor* work correctly, while preserving correct behavior for other applications.

3-SequelPro *Sequel Pro* [4] is a fast, easy-to-use Mac database management application for *MySQL*. It allows user to connect to database with socket or ssh. We experienced the non-responsiveness of *Sequel Pro* when it lost network connection and tried reconections.

The tracing log contains two cases: 1) a quick network connection during login, and 2) *Sequel Pro* lost connection for a while. Although Argus identified the spinning vertex and similar vertex with ease, the backward slicing from similar vertex encountered multiple incoming edges, including one from a kernel thread, where batching processing from different applications happens. Interaction is helpful and reduces the noise in the path. With compariion to the causal path for similar vertex, Argus reveals that the main thread is blocking on a kernel thread, which waits for a ssh thread. Thus we conclude the root cause is the main thread blocking for network IO.

4-Installer *Installer* [5] is an application that extracts and installs files out of *.pkg* packages in macOS. When *Installer*

pops up a window for privileged permission during the installation of *jdk-7u80-macosx-x64*, moving the cursor out of the popup window triggers a spinning cursor.

Argus successfully records the baseline scenario with the following operations. We first type in password in the popup window and then click the back button to reproduce the spinning case by moving cursor. Examining the spinning vertex and its similar vertex, Argus reveals the main thread is waiting for *authd* which blocks on a semaphore. Further diagnosis on *authd* reveals the root cause is the *SecurityAgent*. It processes user input and wakes up *authd* in normal case, but fails to notify *authd* in spinning case. We would suggest an operation for *SecurityAgent* to inform *authd* when the user input box loses cursor, so that the main thread in *Installer* receives proper signal from *authd*.

6.3.2 Long Running

In this section, we discuss the cases where the spinning vertex is busy on the CPU. Most text editing apps fall into this bug category. We studied bugs on *TeXstudio*, *TextEdit*, *Microsoft Word*, *SublimeText*, *TextMate* and *CotEditor*, to reveal the root causes.

5-TeXStudio *TeXstudio* [6] is an integrated writing environment for creating LaTeX documents. We noticed a user reported spinning cursor on the modification of bibliography (bib) file. Although the issue was closed by the developer for insufficient information, we reproduced it with our bib file around 500 items in a *TeXStudio* tab. Each time when we touch the file in other editors like vim, a spinning cursor appears in the application’s window.

Argus recognizes the spinning vertex belongs to the category of LongRunning. The causal path sliced from the spinning vertex by Argus reveals the long-running function is a callback from daemon *fseventd*, and the long processing segment is busy to *CalculateGrowingBlockSize*, even without modifications to the file. The advantage of Argus over other debugging tools is it narrows down the root cause with the inter-processes execution path.

6-TextEdit *TextEdit* is a simple word processing and text editing tool from Apple, which often hangs on editing large files. When Argus is used to diagnose this issue, the heuristics of choosing the most recent edge is powerful enough to get the causal path.

The event graph reveals a communication pattern where a kernel thread is woken from I/O by another kernel thread; the woken kernel thread processed a timer callback function armed by *TextEdit* and finally woke a *TextEdit* thread. In the pattern, the kernel thread has two incoming edges. One is from another kernel thread’s IO completion, and the other is from *TextEdit*’s timer creation. It not hard to reveal the high level semantics. *TextEdit* arms a timer for IO work, and the kernel thread gets the notification for the completion of IO and processes the timer callback. The success of heuristics

BUG-ID	costly API	UI
8-Notes	1)NSDetectScrollDevicesThe -nInvokeOnMainQueue	system defined event
9-SIText	1)px_copy_to_clipboard 2) _CFToUTF8Len	key c
10-TextMate	1)-[OakTextView paste:] 2)CFAttributedStringSet 3)TASCIIEncoder::Encode	key v
11-CotEditor	1)CFStorageGetValueAtIndex 2)-[NSBigMutableString characterAtIndex:]	key Return

Table 4. Root cause of spinning cursor in editing Apps

is not surprising because the most recent edge in the vertex reflects the purpose of the execution segment. Although the heuristics works for the kind of simple application, it is not general enough to succeed for all patterns.

7-MSWord *Microsoft Word* is a large and complex piece of software. Argus can analyze the event graph, but it identifies multiple possible root causes: the length of path interactively sliced from the spinning vertex is 67, while the slicing with heuristics of choosing the most recent edge generates a path of 136 vertices.

We compared the paths and find they diverge from the third vertex backward from spinning vertex. In the vertex, a *Microsoft Word* thread is woken by another *Microsoft Word* thread, and launches a service *NSServiceController-CopyServiceDictionary*. The woken thread sends a message to *launchd* to register the new service and waits for a reply message. With the most recent edge heuristics in automatic slicing, Argus chooses the reply from *launchd* as its predecessor. However, a user can more accurately identify that the execution segment is on behalf of another *Microsoft Word* thread. We rely on user interaction in this case to find the true root cause, since Argus is likely identifies all possibilities without priority.

Other Editing Apps Select, copy, paste, delete, insert and save are common operations for text editing. However, these operations on a large context usually trigger spinning cursors. In Table 4, we list the root causes reported by Argus, including the most costly functions in the event handler, and the user input event (derived from path slicing).

6.3.3 Summary

Overall, in the case of simple text editing applications, Argus can identify the UI event that causes a spinning cursor by merely relying on a few heuristics. However, these heuristics may make the wrong decision in complicated cases, and misidentify the relationships between intra/inter-thread events. It is unlikely that there exists a single graph search method that works in all cases, e.g. when given the choice between multiple incoming edges, the most recent match is sometimes correct, but sometimes not. This is why our

system relies on expert knowledge of users to reconstruct a developer’s intent and accurately diagnose performance issues.

6.4 Performance Evaluation

In this section we present the performance impact of the live deployment of Argus. We deployed Argus on a MacBookPro9,2, which has Intel Core i5-3210M CPU with 2 cores and 4 thread, 10GB DDR3 memory and a 1T SSD.

Argus has a very small space overhead with the configuration of its tracing tool. It uses the ring buffer with configured 2G by default to collect tracing events. The memory used to store events is fixed to 512M by Apple, which is pretty low with regards to the memory usage of modern applications. In the remaining of this section, we measure Argus’s overhead overall with iBench scores, IO throughput degradation with bonnie++, iotzone and CPU overhead with chromium benchmarks.

iBench We first show the five runs of iBench with and without Argus to evaluate the overall performance. The machine is clean booted for each run, and the higher score means it performs better. As shown in Table 6, their performance are almost of no difference, only 0.13% degradation on average.

IO Throughput Next, we evaluate the IO throughput with bonnie++ and iotzone. As shown in the Table 7, the throughputs of sequential read and write by characters with and without Argus are almost same. Read and write by block impose less than 10% overhead in both microbenchmarks, bonni++ and iotzone. With selected events in our system, the tracing tool integrated in Argus only adds 5% IO overhead on average.

CPU We evaluate Argus’s CPU overhead with chromium benchmarks by recording their time usage on real, user and sys. Although the overhead on sys is relatively higher than other two, due to the tracing events usually crossing the kernel boundary, they are not triggered too frequently in our daily software usage, including browsers. The time cost is mostly under 5%, except the *dummy_benchmark.histogram*. As shown in Table 5, the time overhead for real, user and sys are 7%, 5% and 40% respectively.

7 Related Work

While there is currently no system that can help users debug performance issues in closed-source applications on proprietary macOS, several active research topics are closely related.

Event tracing. Panappticon [32] monitors a mobile system and uses the trace to characterize the user transactions of mobile apps. Although it aims to track system-wide events and correlate them without developer input, it supports only

Chromium Benchmark (in seconds)	with Argus			without Argus			Overhead		
	real	user	sys	real	user	sys	real	user	sys
system_health.memory_desktop	11592	18424	1821	11317	18401	1415	0.02	0.00	0.29
rasterize_and_record_micro.top_25	1579	2142	135	1654	2166	116	-0.05	-0.01	0.16
blink_perf	16210	17227	959	15877	16724	766	0.02	0.03	0.25
webrtc	726	2023	225	725	2130	168	0.00	-0.05	0.34
memory.desktop	1231	2238	267	1188	2200	190	0.04	0.02	0.41
loading.desktop.network_service	24580	52751	6294	23696	52327	4197	0.04	0.01	0.50
dromaeo	206	227	15	192	212	12	0.07	0.07	0.29
dummy_benchmark.histogram	49	48	8	33	36	4	0.50	0.32	0.96
v8.browsing_desktop	2462	4489	491	2325	4440	303	0.06	0.01	0.62
octan.desktop	112	142	8	98	124	5	0.14	0.15	0.44
speedometer	618	802	31	600	782	24	0.03	0.03	0.32
page_cycler_v2.typical_2	8020	14435	1453	7847	14215	1019	0.02	0.02	0.43
smoothness.oop_rasterization.top_25_smooth	864	1450	156	833	1412	126	0.04	0.03	0.24
AVERAGE	-			-			0.07	0.05	0.4

Table 5. Chromium benchmark

	1st	2nd	3rd	4th	5th
without Argus	5.98	6.23	6.18	6.05	6.28
with Argus	6.29	6.01	6.09	6.28	6.01
average overhead	0.13%				

Table 6. Score From iBench

	kb/s	With Argus	Without Argus	overhead
bonnie++	read char	21922	22149	0.01
	sequential read block	226931	244089	0.07
	rewrite	246807	267491	0.08
	write char	22924	22936	0.00
	write block	4073361	4396387	0.07
seq	file create	17391	17381	0.00
	file delete	18089	19401	0.07
random	create	17472	17887	0.02
	delete	8849	9567	0.08
iozone	initial write	1199453	1318572	0.09
	rewrite	3663066	4059912	0.10
	average	-	-	0.05

Table 7. IO throughput with bonnie++ and iozone

two models of communication: work queue and thread pooling. AppInsight [25] instruments application to identify the critical execution path in a user transaction. It supports the event callback pattern, and does not trace across process or app boundaries. Magpie [12] monitors server applications in Windows with the goal to model the normal behaviors of a server application in response to a workload. This model further helps detecting anomalies statistically. Magpie requires a manual-written event schema for all involved applications to capture precise request graphs, whereas Argus has a simple, application-agnostic schema for system-wide tracing and enables users to provide more application-specific knowledge on demand.

Aguilela [10] uses timing analysis to correlate messages to recover their input-output relations while treating the application as a black box. XTrace, Pinpoint and etc [14, 15, 19] trace the path of a request through a system using a unique identifier attached to each request and stitch traces together with the identifier. Argus comes up violation patterns and does not assume the presence of a unified identifier in closed-source, third-party applications, frameworks, and libraries.

Performance anomaly detection. Several systems detect performance anomalies automatically. [20, 31] leverage the user logs and call stacks to identify the performance anomaly. [16, 17, 27, 29] apply the machine learning method to identify the unusual event sequence as an anomaly. [30] generates the wait and waken graph from sampled call stacks to study a case of performance anomaly.

These systems are orthogonal to Argus as Argus’s goal is to diagnose an already-detected performance anomaly. These systems can help Argus by detecting more accurately when a performance issue arises.

8 Conclusion

Our key insight in this paper is that causal tracing is inherently inaccurate. The inaccuracy unlikely gets remedied by heuristics so as to generate a feasible request graph for automatic diagnosis on performance issues. We built Argus, a practical system for effectively debugging performance issues despite inaccurate causal tracing. It lets a user provide domain knowledge on demand to mitigate the inaccuracy. Compared to all upfront knowledge, this method is more general and efficient given the various programing paradigms.

References

- [1] https://opensource.apple.com/source/system_cmds/system_cmds-671.10.3/trace.tproj.
- [2] <https://opensource.apple.com/source/xnu>.
- [3] <http://mirror.informatimago.com/next/developer.apple.com/releasesnotes/DeveloperTools/TwoLevelNamespaces.html>.
- [4] <https://www.sequelpro.com>.
- [5] [https://en.wikipedia.org/wiki/Installer_\(macOS\)](https://en.wikipedia.org/wiki/Installer_(macOS)).
- [6] <https://www.textstudio.org>.
- [7] The Chromium Projects. <https://www.chromium.org>, 2008.
- [8] Issue 115920: Response time can be really long with some IMEs (e.g. Pinyin IME (Apple), Sogou Pinyin IME). <https://bugs.chromium.org/p/chromium/issues/detail?id=115920>, 2012.
- [9] <https://github.com/Eun/DisableMonitor>, 2018.
- [10] Aguilera, Marcos K and Mogul, Jeffrey C and Wiener, Janet L and Reynolds, Patrick and Muthitacharoen, Athicha. Performance debugging for distributed systems of black boxes. In *ACM SIGOPS Operating Systems Review*, 2003.
- [11] Attariyan, Mona and Chow, Michael and Flinn, Jason. X-ray: Automating root-cause diagnosis of performance anomalies in production software. In *Presented as part of the 10th USENIX Symposium on Operating Systems Design and Implementation OSDI 12*, 2012.
- [12] Barham, Paul and Donnelly, Austin and Isaacs, Rebecca and Mortier, Richard. Using Magpie for request extraction and workload modelling. In *OSDI*, 2004.
- [13] Benjamin H. Sigelman and Luiz Andr   Barroso and Mike Burrows and Pat Stephenson and Manoj Plakal and Donald Beaver and Saul Jaspam and Chandan Shanbhag. Dapper, a Large-Scale Distributed Systems Tracing Infrastructure. Technical report, 2010.
- [14] Chen, Mike Y and Kiciman, Emre and Fratkin, Eugene and Fox, Armando and Brewer, Eric. Pinpoint: Problem determination in large, dynamic internet services. In *Proceedings International Conference on Dependable Systems and Networks*, 2002.
- [15] Chow, Michael and Meisner, David and Flinn, Jason and Peek, Daniel and Wensch, Thomas F. The mystery machine: End-to-end performance analysis of large-scale internet services. In *11th USENIX Symposium on Operating Systems Design and Implementation OSDI 14*, 2014.
- [16] Cohen, Ira and Chase, Jeffrey S and Goldszmidt, Moises and Kelly, Terence and Symons, Julie. Correlating Instrumentation Data to System States: A Building Block for Automated Diagnosis and Control. In *OSDI*, 2004.
- [17] Du, Min and Li, Feifei and Zheng, Guineng and Srikumar, Vivek. Deeplog: Anomaly detection and diagnosis from system logs through deep learning. In *Proceedings of the 2017 ACM SIGSAC Conference on Computer and Communications Security*, 2017.
- [18] Dunlap, George W and Lucchetti, Dominic G and Fetterman, Michael A and Chen, Peter M. Execution replay of multiprocessor virtual machines. In *Proceedings of the fourth ACM SIGPLAN/SIGOPS international conference on Virtual execution environments*, 2008.
- [19] Fonseca, Rodrigo and Porter, George and Katz, Randy H and Shenker, Scott and Stoica, Ion. X-trace: A pervasive network tracing framework. In *Proceedings of the 4th USENIX conference on Networked systems design & implementation*, 2007.
- [20] Han, Shi and Dang, Yingnong and Ge, Song and Zhang, Dongmei and Xie, Tao. Performance debugging in the large via mining millions of stack traces. In *Proceedings of the 34th International Conference on Software Engineering*, 2012.
- [21] Harter, Tyler and Dragga, Chris and Vaughn, Michael and Arpaci-Dusseau, Andrea C and Arpaci-Dusseau, Remzi H. A file is not a file: understanding the I/O behavior of Apple desktop applications. *ACM Transactions on Computer Systems (TOCS)*, 2012.
- [22] Galen Hunt and Doug Brubacher. Detours: Binary interception of win32 functions. In *3rd unix windows nt symposium*, 1999.
- [23] King, Samuel T and Dunlap, George W and Chen, Peter M. Debugging operating systems with time-traveling virtual machines. In *Proceedings of the annual conference on USENIX Annual Technical Conference*, 2005.
- [24] Nagaraj, Karthik and Killian, Charles and Neville, Jennifer. Structured comparative analysis of systems logs to diagnose performance problems. In *Proceedings of the 9th USENIX conference on Networked Systems Design and Implementation*, 2012.
- [25] Ravindranath, Lenin and Padhye, Jitendra and Agarwal, Sharad and Mahajan, Ratul and Obermiller, Ian and Shayandeh, Shahin. Applnsight: mobile app performance monitoring in the wild. In *Presented as part of the 10th USENIX Symposium on Operating Systems Design and Implementation OSDI 12*, 2012.
- [26] Reynolds, Patrick and Killian, Charles Edwin and Wiener, Janet L and Mogul, Jeffrey C and Shah, Mehul A and Vahdat, Amin. Pip: Detecting the Unexpected in Distributed Systems. In *NSDI*, 2006.
- [27] Saidi, Ali G and Binkert, Nathan L and Reinhardt, Steven K and Mudge, Trevor. Full-system critical path analysis. In *ISPASS 2008-IEEE International Symposium on Performance Analysis of Systems and Software*, 2008.
- [28] Xiong, Weiwei and Park, Soyeon and Zhang, Jiaqi and Zhou, Yuanyuan and Ma, Zhiqiang. Ad Hoc Synchronization Considered Harmful. In *OSDI*, 2010.
- [29] Xu, Wei and Huang, Ling and Fox, Armando and Patterson, David and Jordan, Michael I. Detecting large-scale system problems by mining console logs. In *Proceedings of the ACM SIGOPS 22nd symposium on Operating systems principles*, 2009.
- [30] Yu, Xiao and Han, Shi and Zhang, Dongmei and Xie, Tao. Comprehending performance from real-world execution traces: A device-driver case. In *ACM SIGPLAN Notices*, 2014.
- [31] Yuan, Ding and Park, Soyeon and Huang, Peng and Liu, Yang and Lee, Michael M and Tang, Xiaoming and Zhou, Yuanyuan and Savage, Stefan. Be conservative: enhancing failure diagnosis with proactive logging. In *Presented as part of the 10th USENIX Symposium on Operating Systems Design and Implementation OSDI 12*, 2012.
- [32] Zhang, Lide and Bild, David R and Dick, Robert P and Mao, Z Morley and Dinda, Peter. Panappticon: event-based tracing to measure mobile application and platform performance. In *2013 International Conference on Hardware/Software Codesign and System Synthesis (CODES+ ISSS)*, 2013.
- [33] Zhao, Xu and Rodrigues, Kirk and Luo, Yu and Yuan, Ding and Stumm, Michael. Non-intrusive performance profiling for entire software stacks based on the flow reconstruction principle. In *12th USENIX Symposium on Operating Systems Design and Implementation OSDI 16*, 2016.